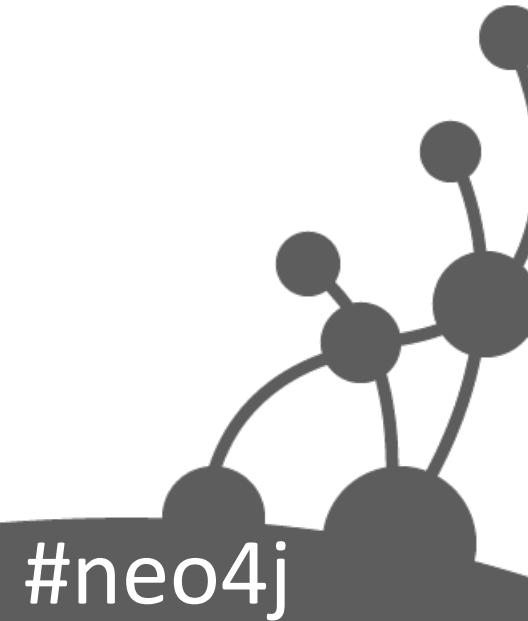




Thinking in Graphs

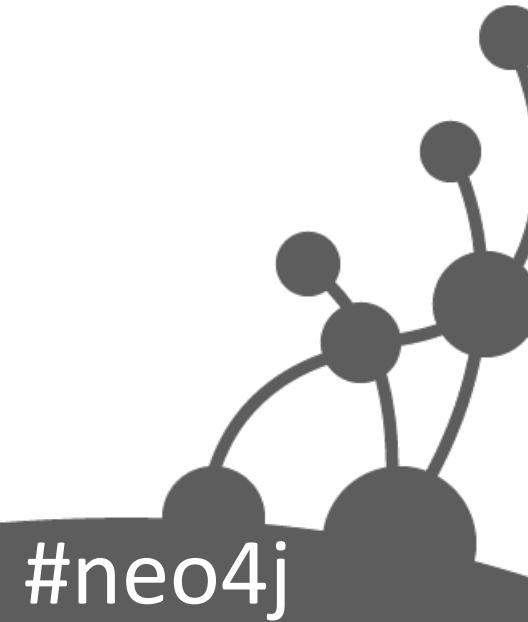
@iansrobinson
ian.robinson@neotechnology.com

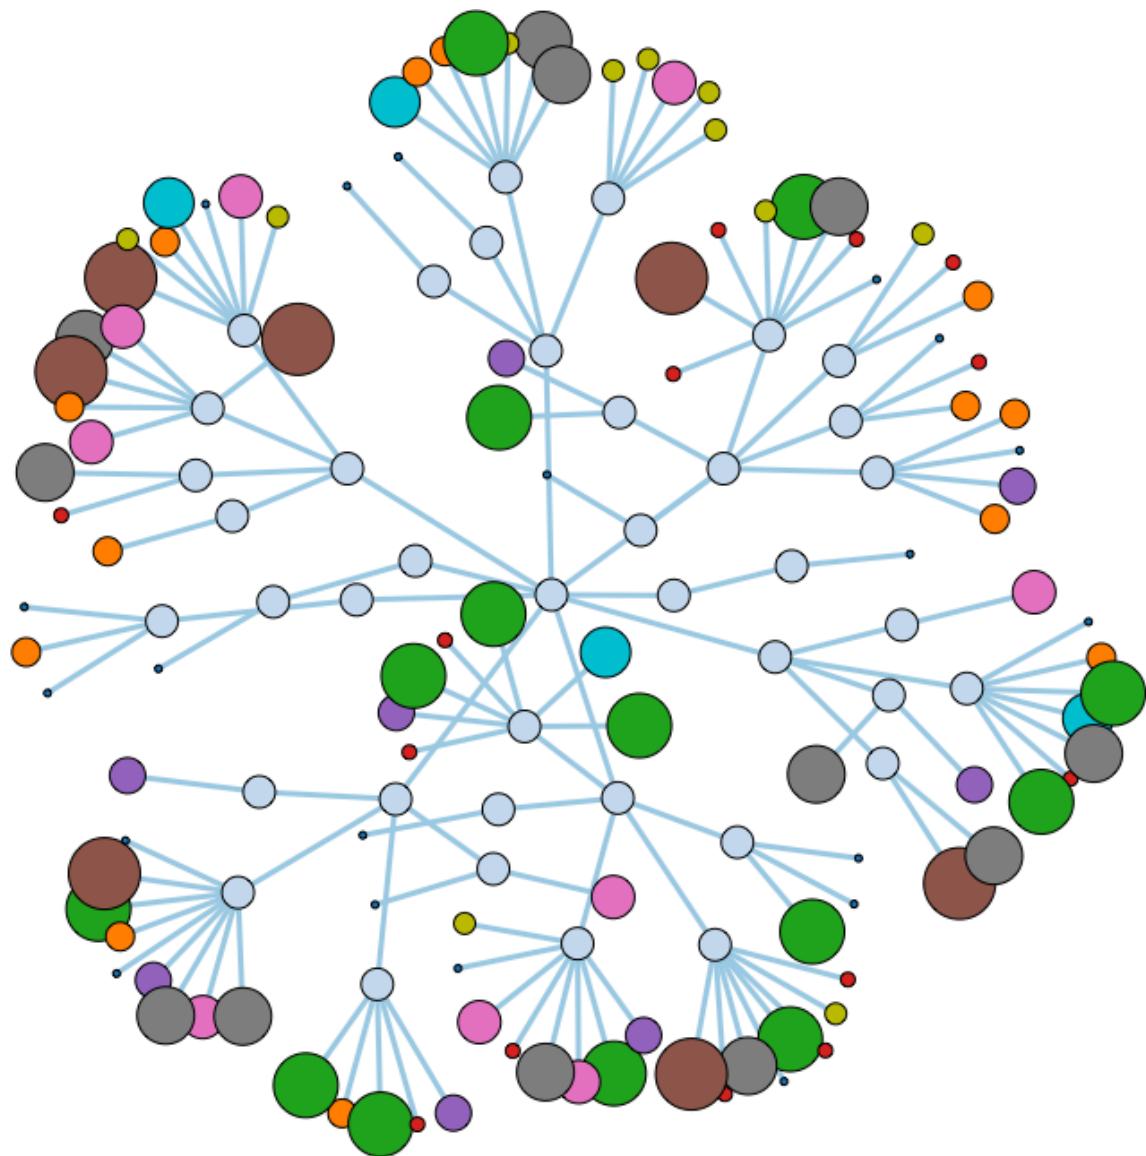


#neo4j

Data Complexity

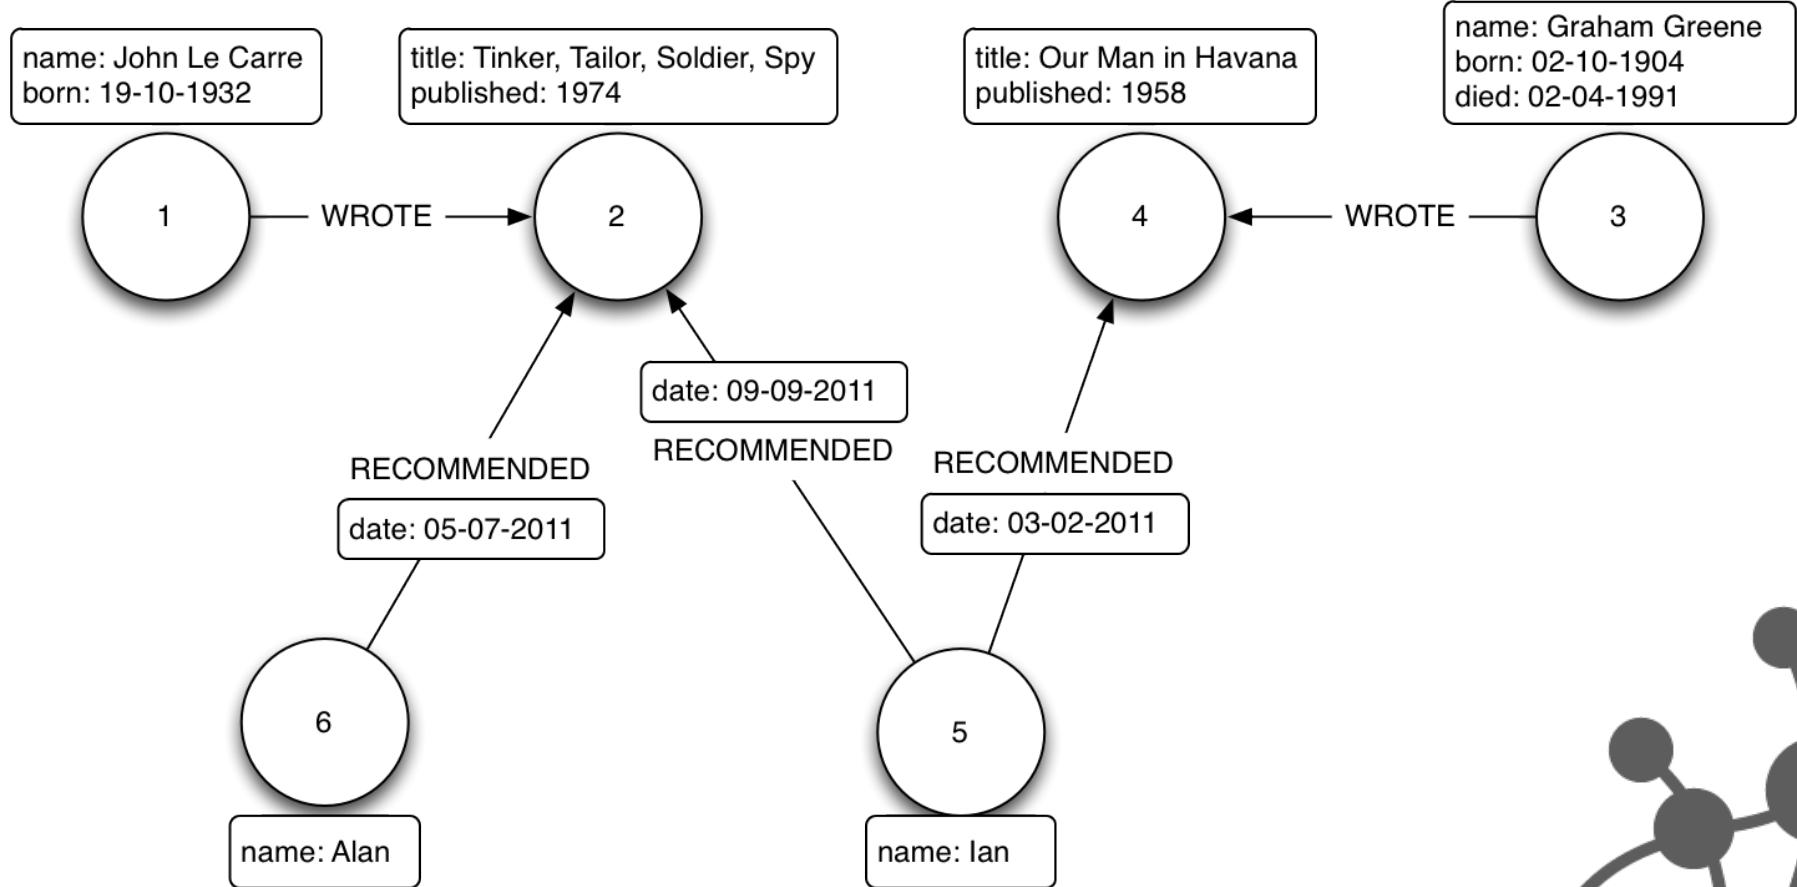
$complexity = f(size, connectedness, semi-structure)$



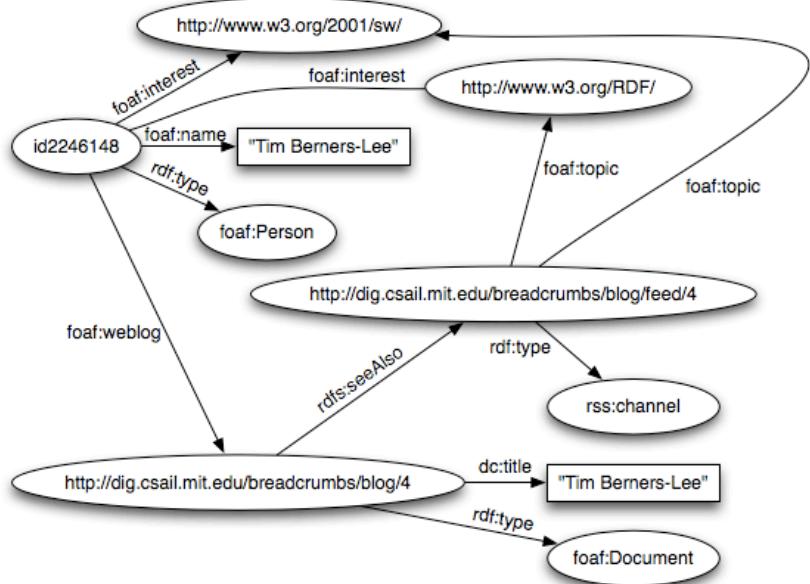
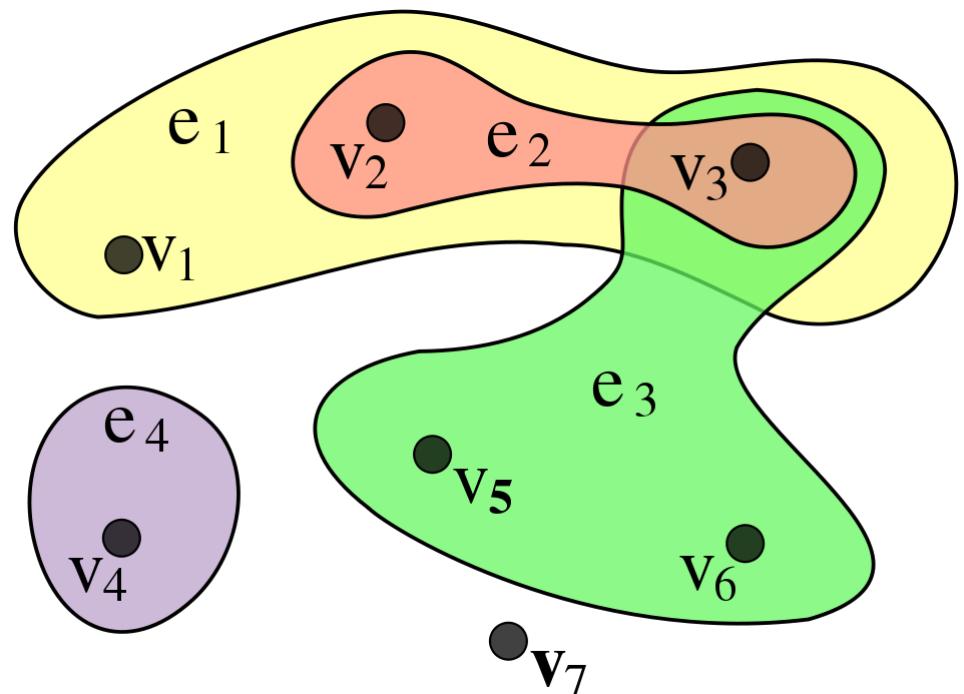


#neo4j

Property Graph Data Model



#neo4j



#neo4j

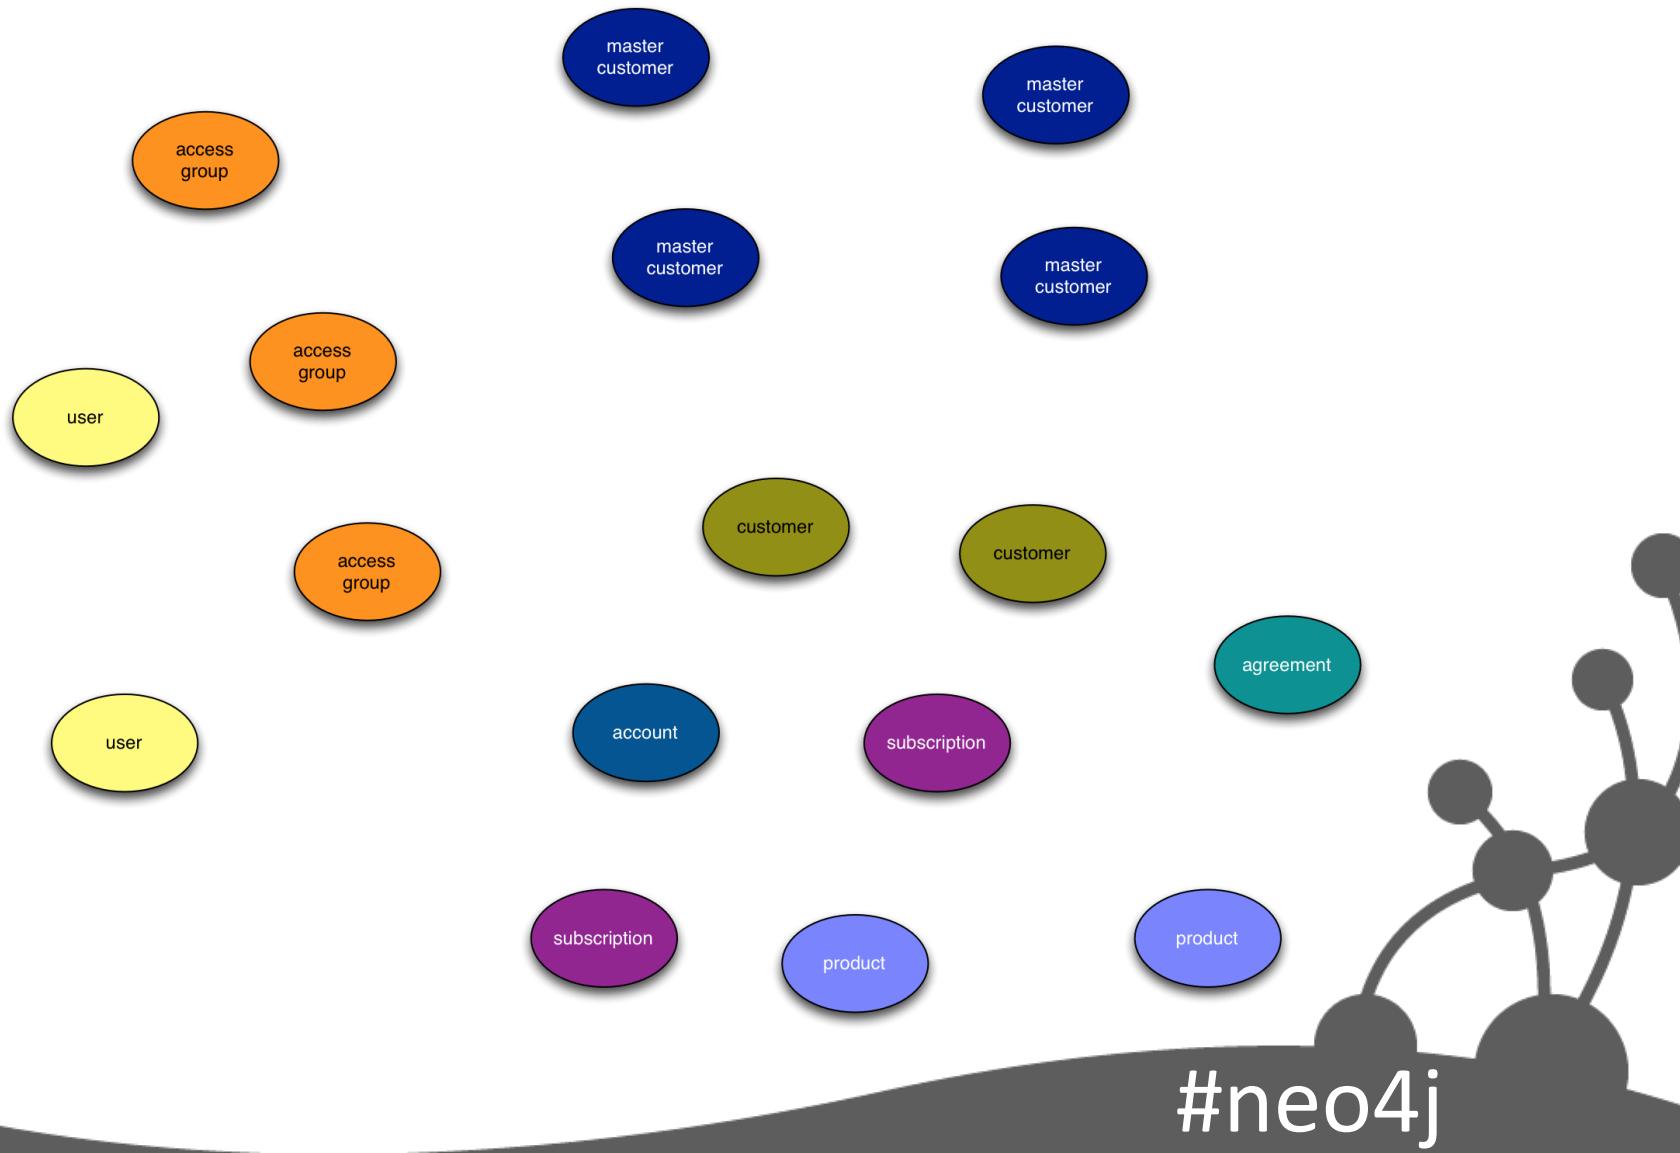
Expressive Graph Data Model

- Complex, densely-connected domains
 - Lots of join tables? Relationships
 - Lots of sparse tables? Semi-structure
- Messy data
 - Ad hoc exceptions
- Relationships as first-class elements
 - Semantic clarity: named, directed
 - Not simply constraints

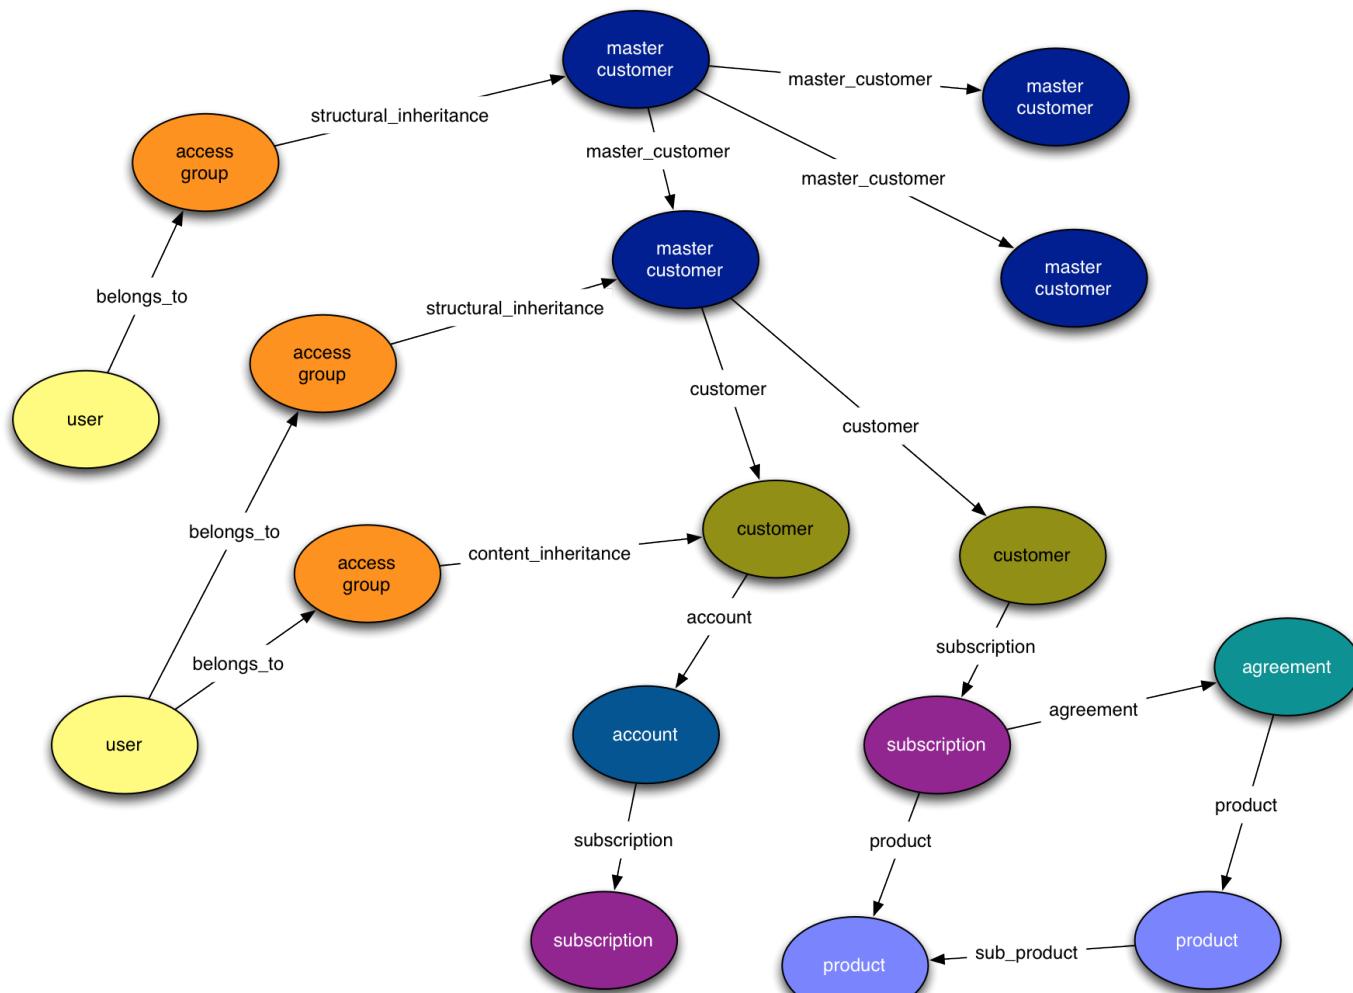


#neo4j

Schema-Free...

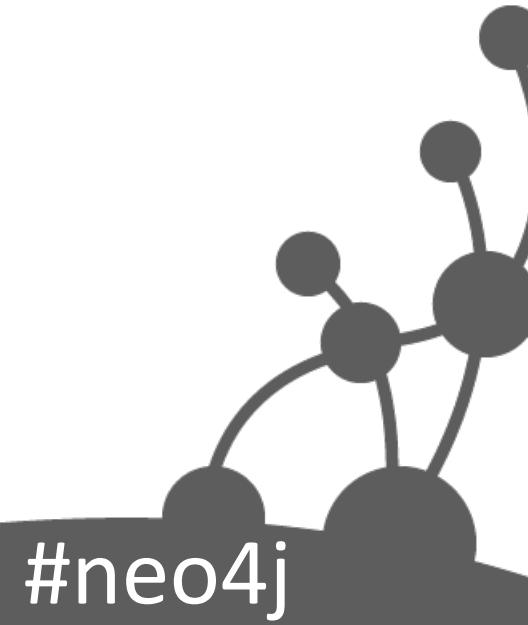
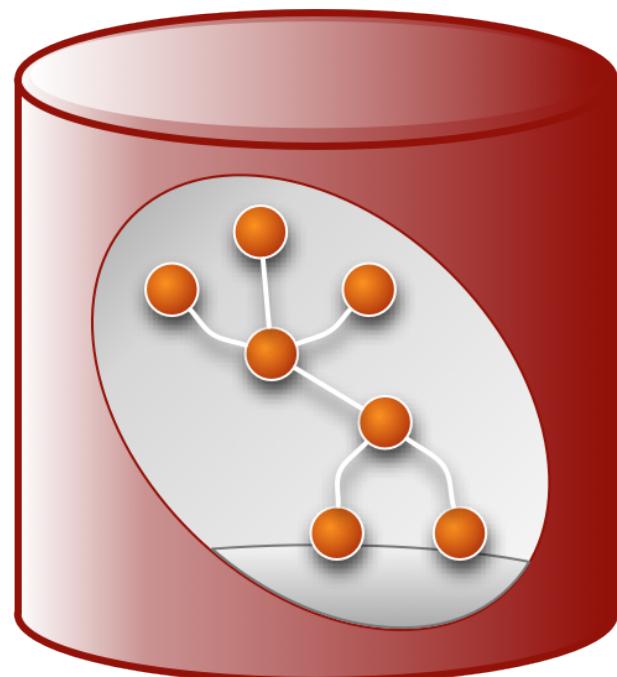


But Structured



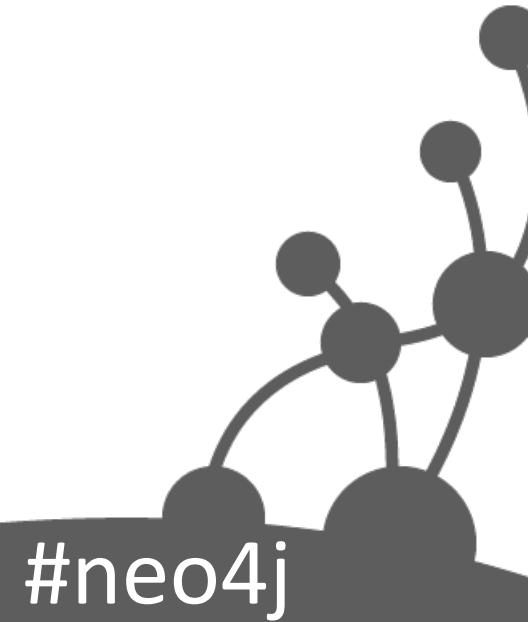
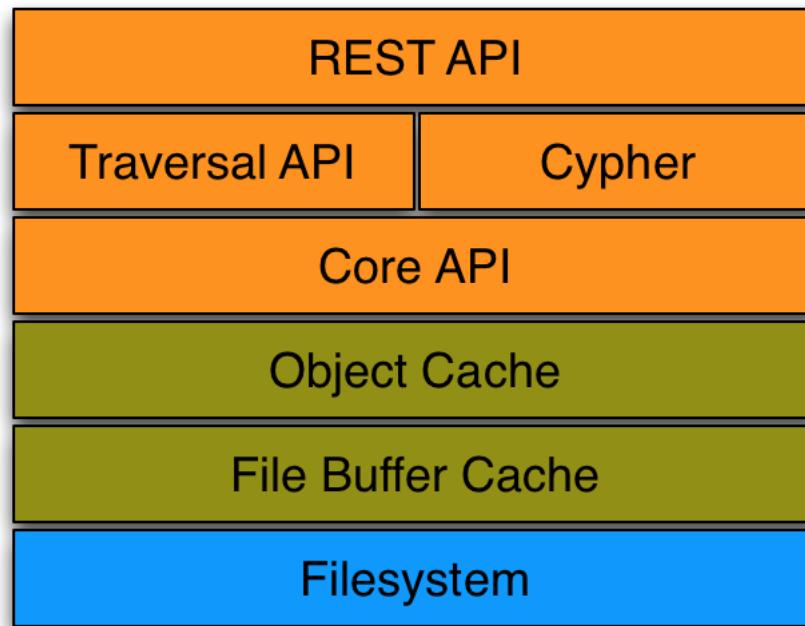
#neo4j

Neo4j



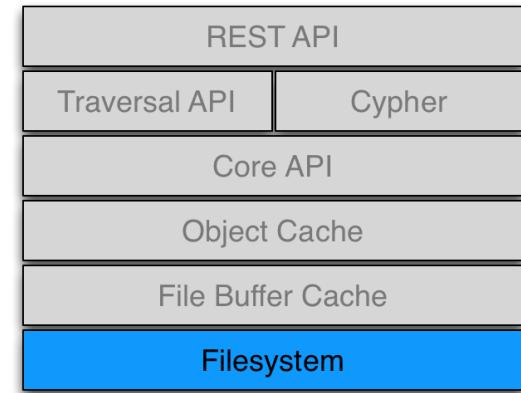
#neo4j

Neo4j Architecture



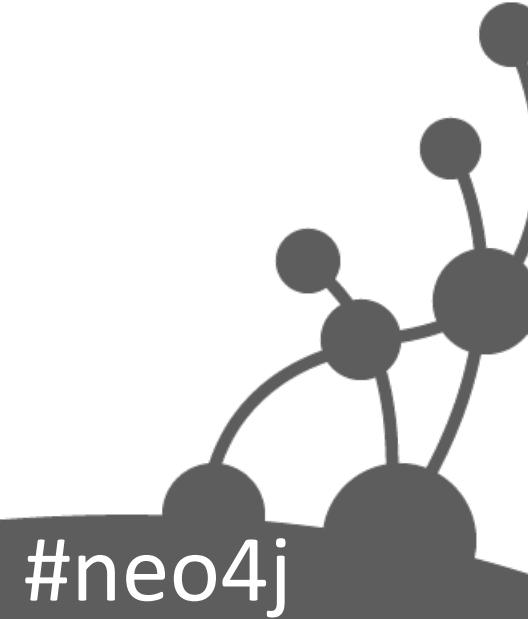
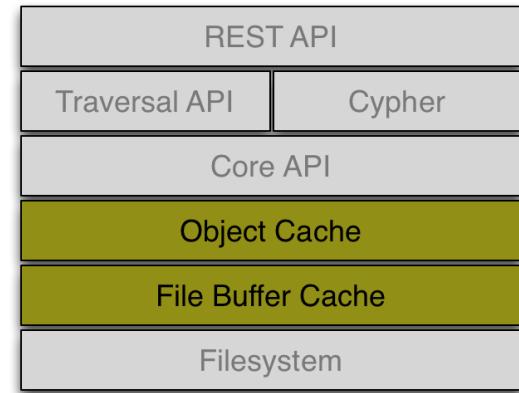
Filesystem

- Neo4j uses several stores:
 - Node
 - Relationship
 - Property
 - including short strings and small arrays
 - String
 - for long strings
 - Array
 - for large arrays



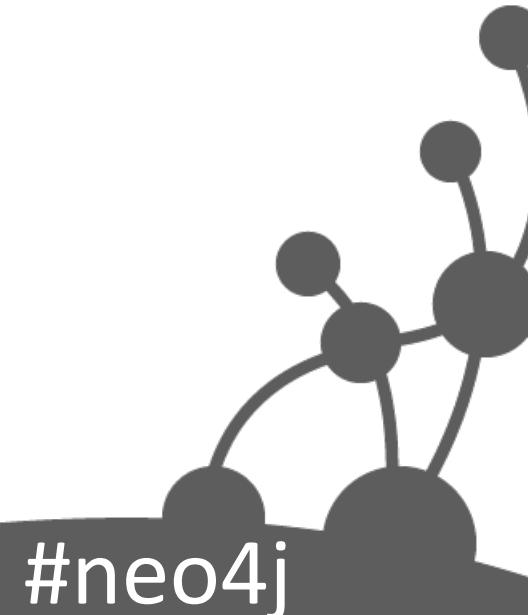
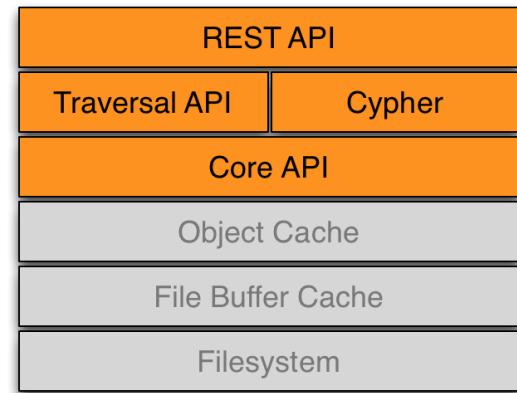
Caches

- File Buffer Cache:
 - Memory-mapped NIO
- Object Cache:
 - Nodes
 - Relationships

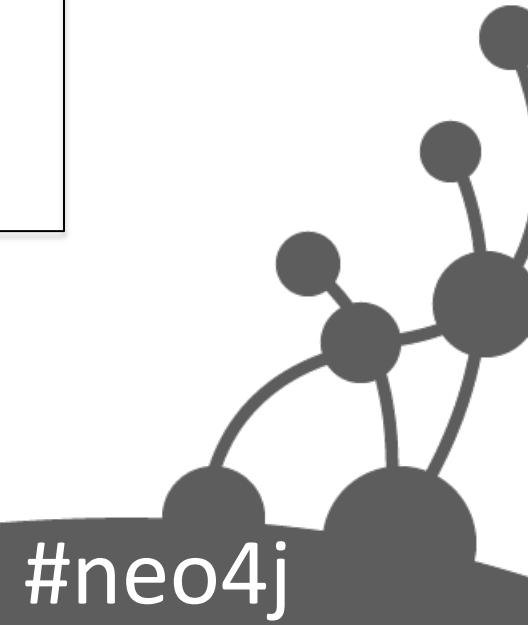
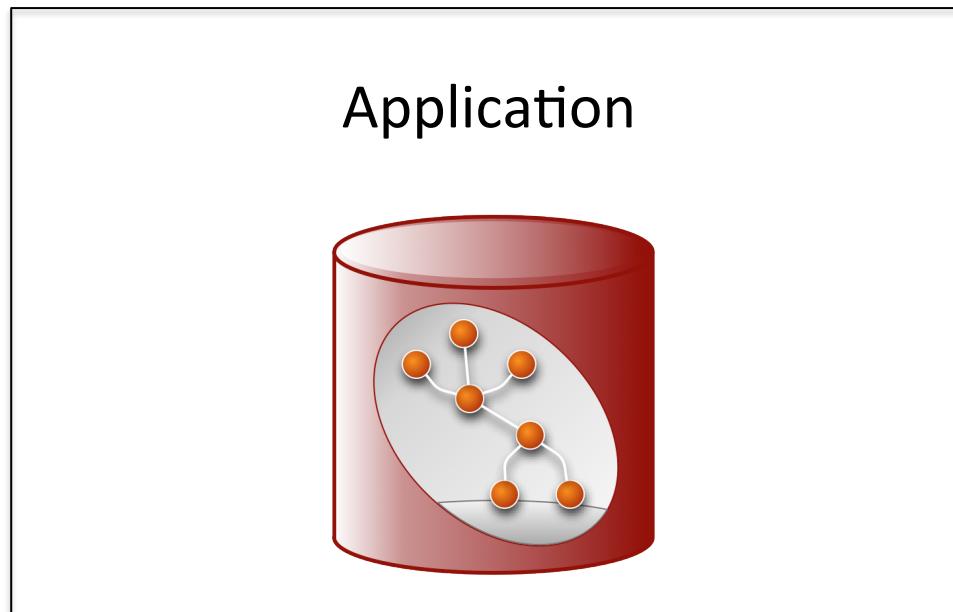


APIs

- Core API:
 - Nodes
 - Relationships
- Traversal API:
 - Lazily spider the graph
- Cypher
 - Graph pattern matching
- REST
 - JSON over HTTP

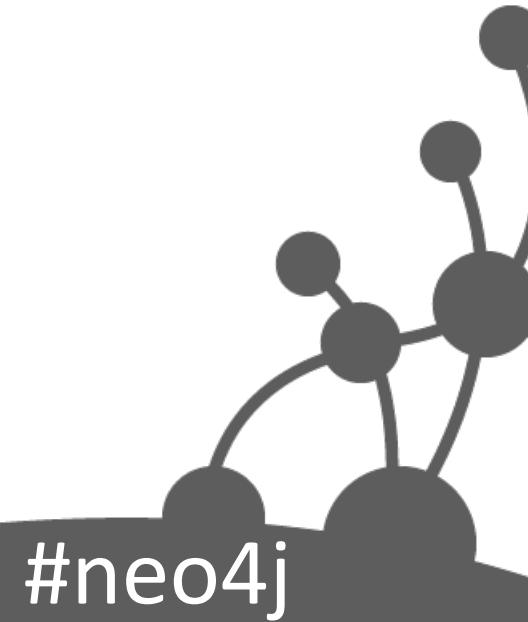
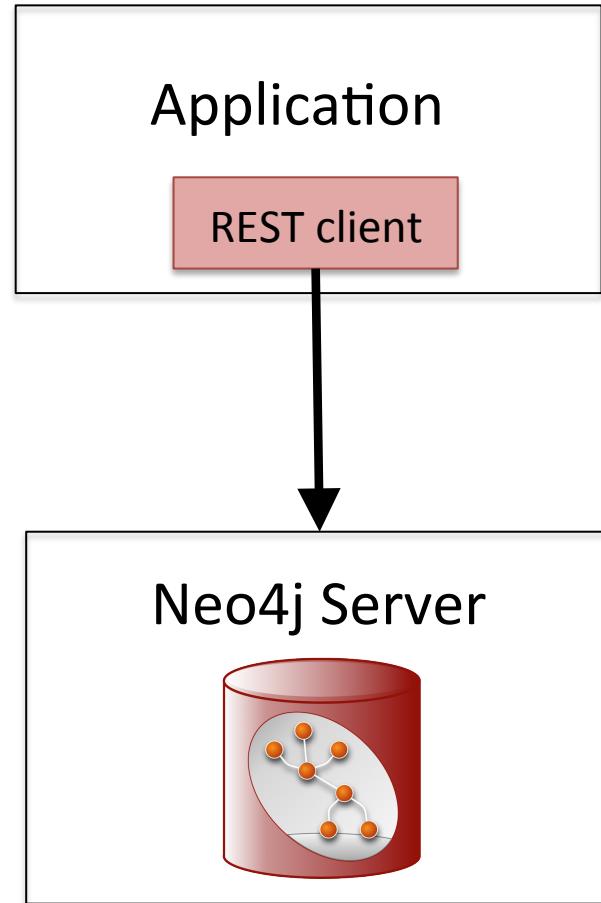


Embedded

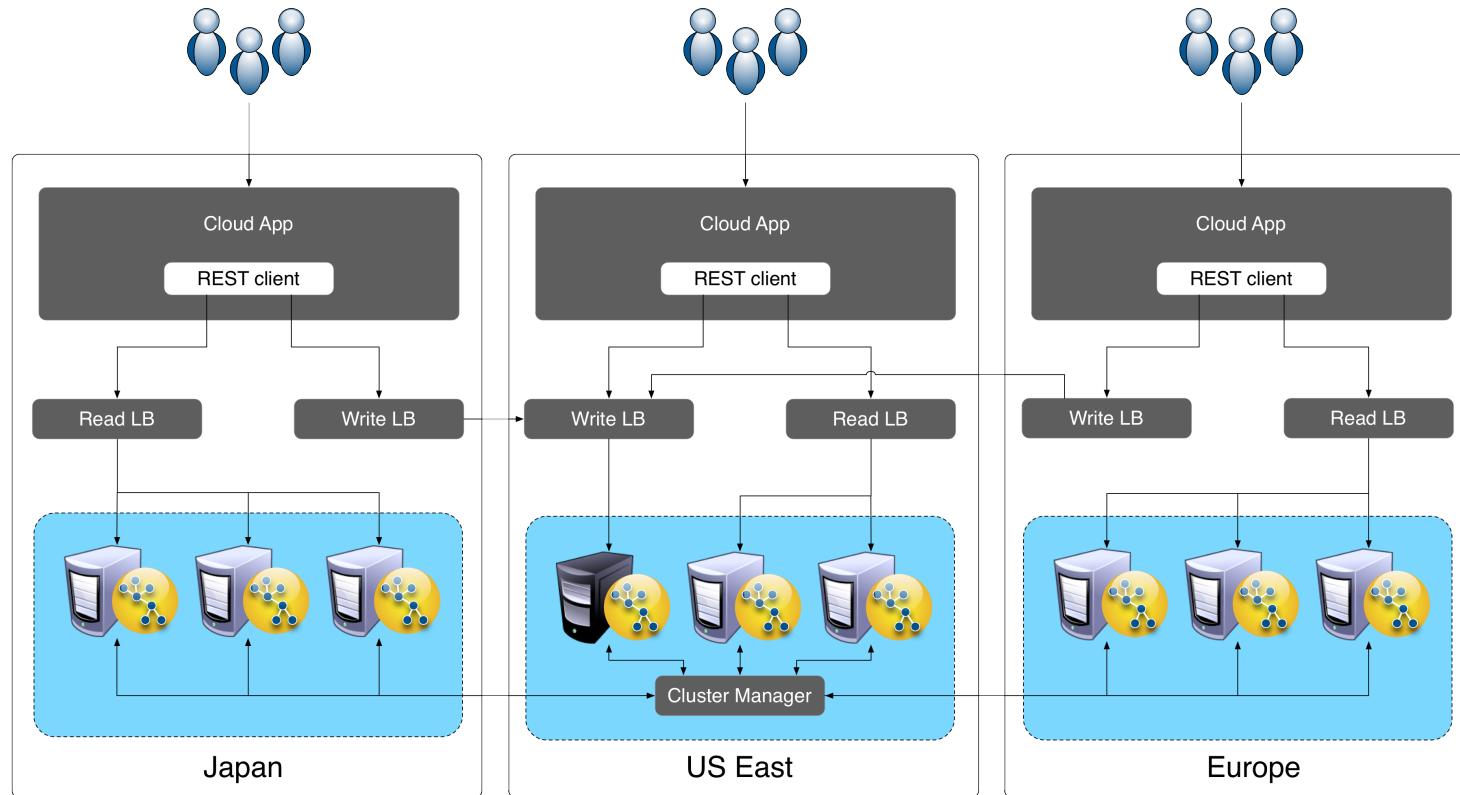


#neo4j

Server

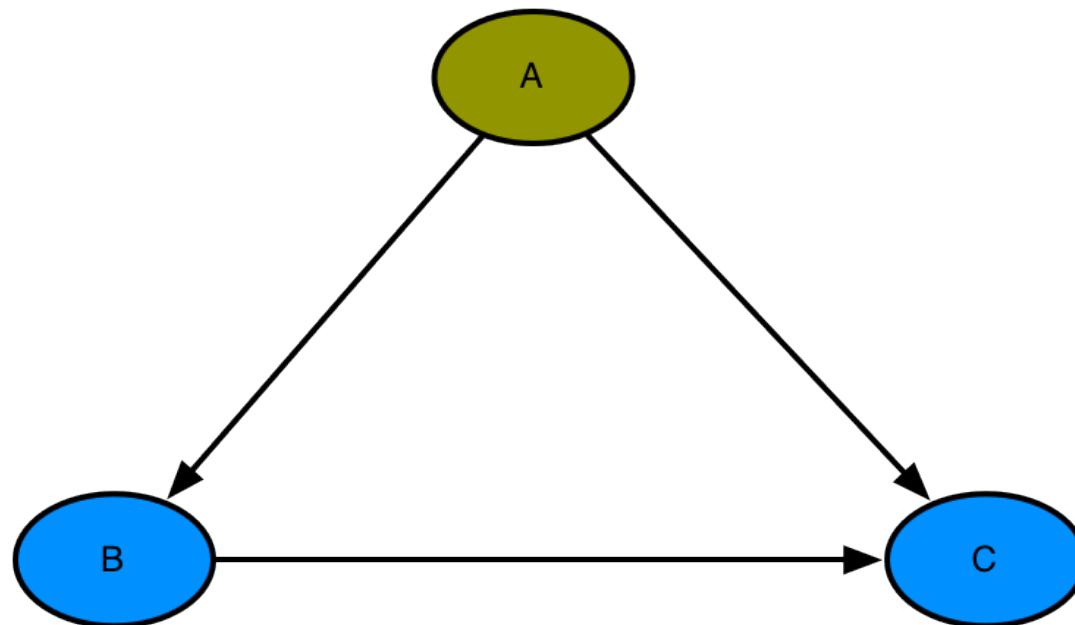


High Availability



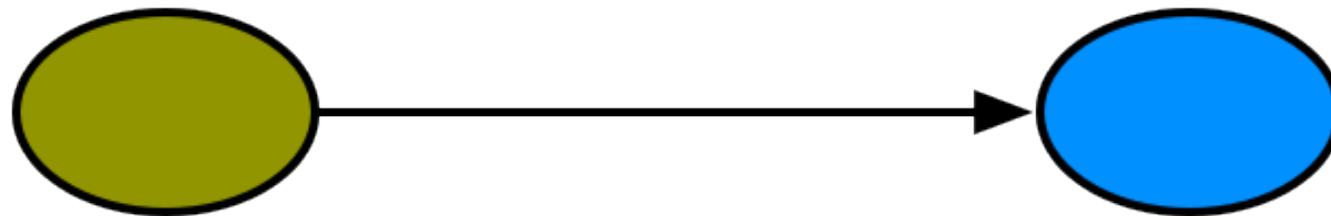
#neo4j

Cypher – A Graph Language

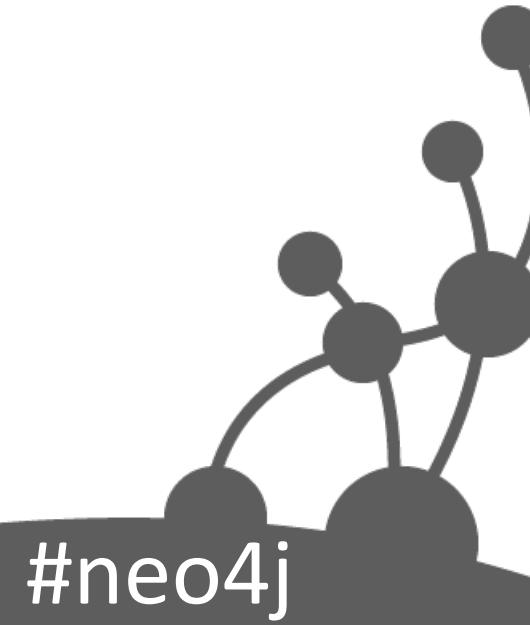


#neo4j

Anonymous Nodes & Rel



○ - -> ○

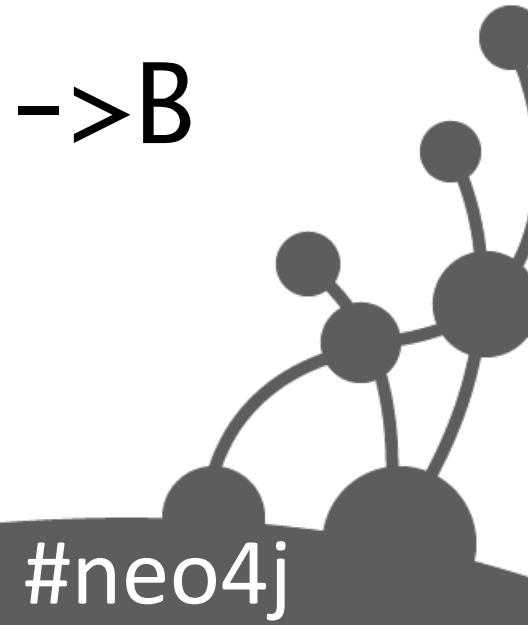


#neo4j

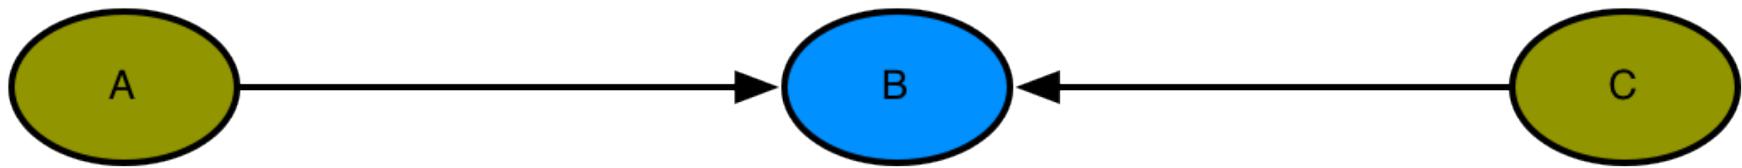
ASCII Art Patterns



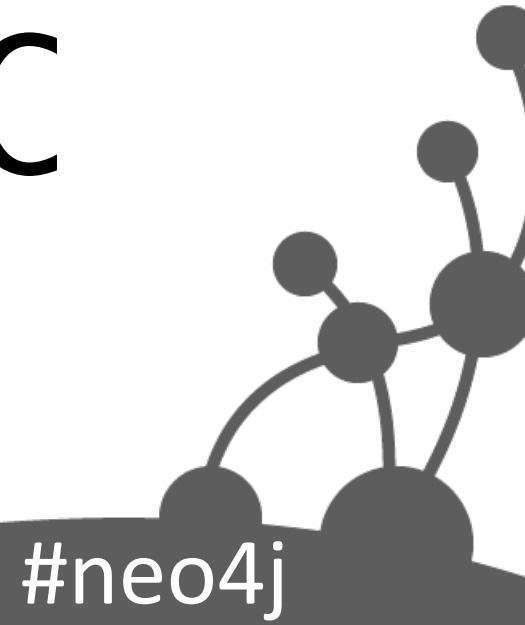
A-[:CONNECTED_TO]->B



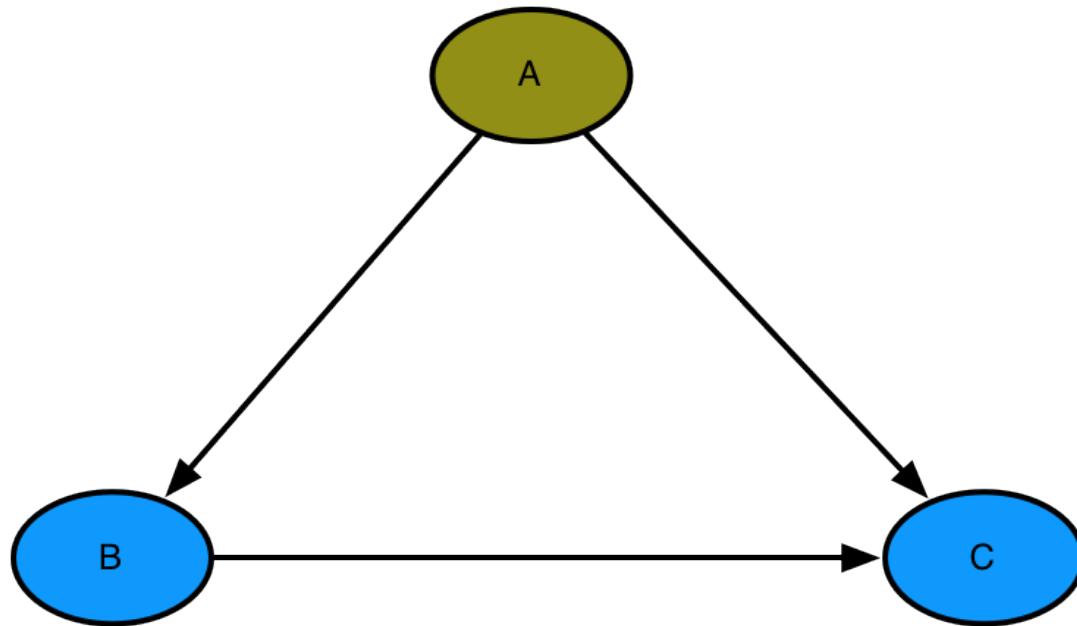
ASCII Art Patterns



A --> B <-- C

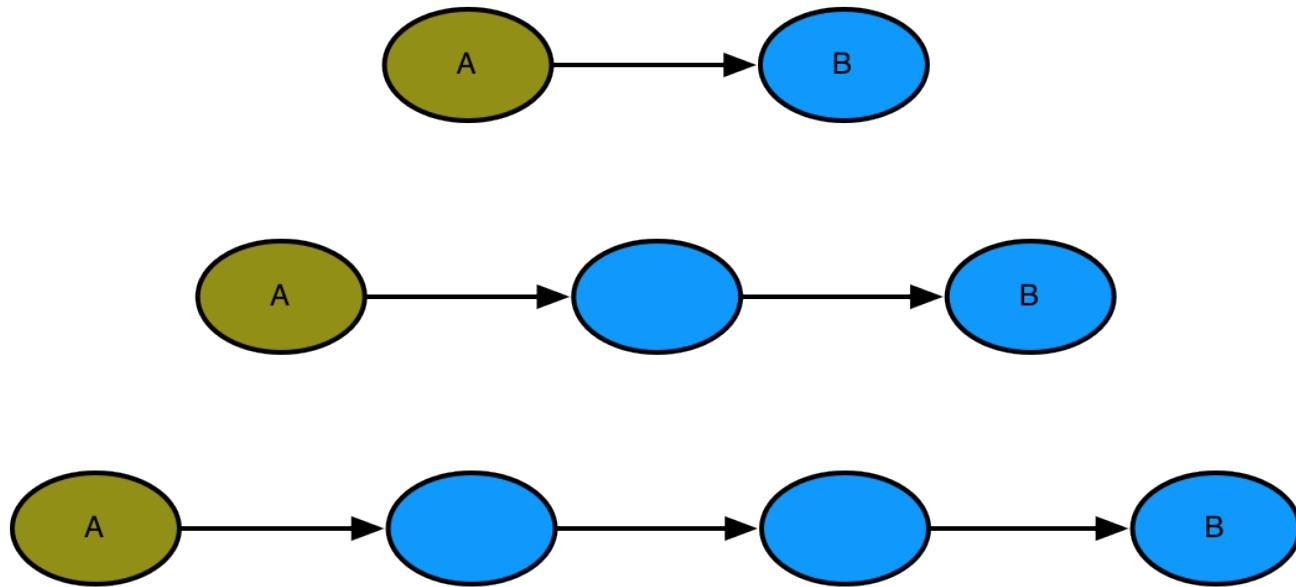


ASCII Art Patterns

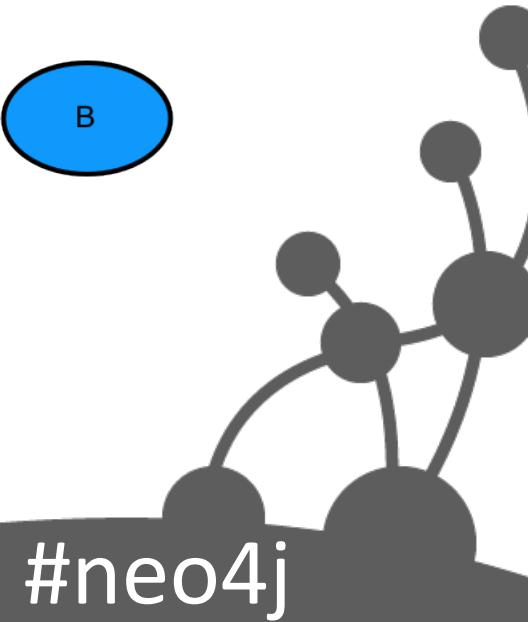


A-->B-->C, A-->C
A-->B-->C<---A

Variable Length Paths



A - [*] -> B



#neo4j

Create Some Data

CREATE

```
(cc {name:'Cost Centre 1'}),  
(lb {name:'Ledger Book 1'}),  
cc-[ :HAS_LEDGER_BOOK ]->lb
```

Add Data

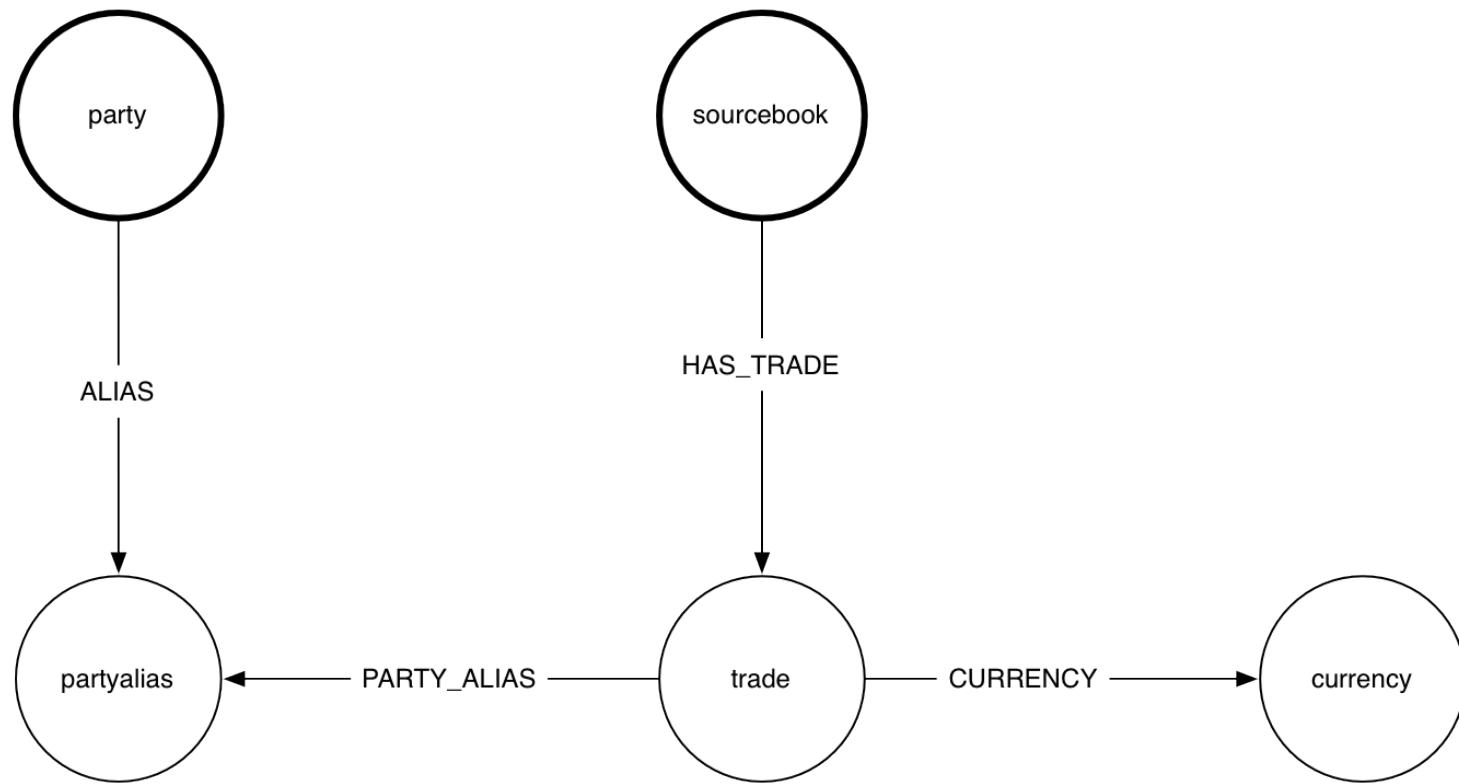
```
START      party=node:party(name='Party 1'),  
          sourcebook=node:source_book(name='Source Book 1')  
  
CREATE UNIQUE sourcebook-[:HAS_TRADE]->(trade {id:'0001'}),  
               trade-[:PARTY_ALIAS]->  
                  (partyalias {name:'Party Alias 1'}),  
               trade-[:CURRENCY]->  
                  (currency {name:'GBP', _label:'currency'}),  
               party-[:ALIAS]->partyalias  
  
RETURN     trade
```

Find Start Nodes in Indexes

```
START      party=node:party(name='Party 1'),  
           sourcebook=node:source_book(name='Source Book 1')  
  
CREATE UNIQUE sourcebook-[:HAS_TRADE]->(trade {id:'0001'}),  
               trade-[:PARTY_ALIAS]->  
                   (partyalias {name:'Party Alias 1'}),  
               trade-[:CURRENCY]->  
                   (currency {name:'GBP', _label:'currency'}),  
               party-[:ALIAS]->partyalias  
  
RETURN     trade
```

Describe New Data

```
START      party=node:party(name='Party 1'),  
          sourcebook=node:source_book(name='Source Book 1')  
CREATE UNIQUE sourcebook-[:HAS_TRADE]->(trade {id:'0001'}),  
                    trade-[:PARTY_ALIAS]->  
                        (partyalias {name:'Party Alias 1'}),  
                    trade-[:CURRENCY]->  
                        (currency {name:'GBP', _label:'currency'}),  
                    party-[:ALIAS]->partyalias  
RETURN    trade
```



#neo4j

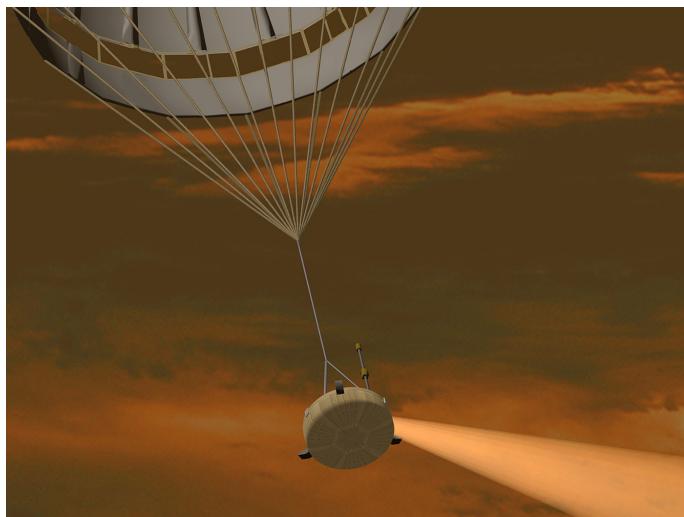
Return New Node

```
START      party=node:party(name='Party 1'),  
          sourcebook=node:source_book(name='Source Book 1')  
  
CREATE UNIQUE sourcebook-[:HAS_TRADE]->(trade {id:'0001'}),  
               trade-[:PARTY_ALIAS]->  
                  (partyalias {name:'Party Alias 1'}),  
               trade-[:CURRENCY]->  
                  (currency {name:'GBP', _label:'currency'}),  
               party-[:ALIAS]->partyalias  
  
RETURN     trade
```

Querying the Graph

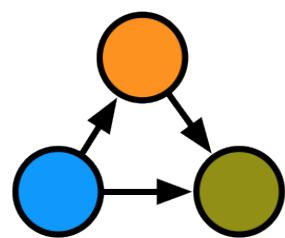
Graph local:

- One or more start nodes
- Explore surrounding graph
- Millions of joins per second

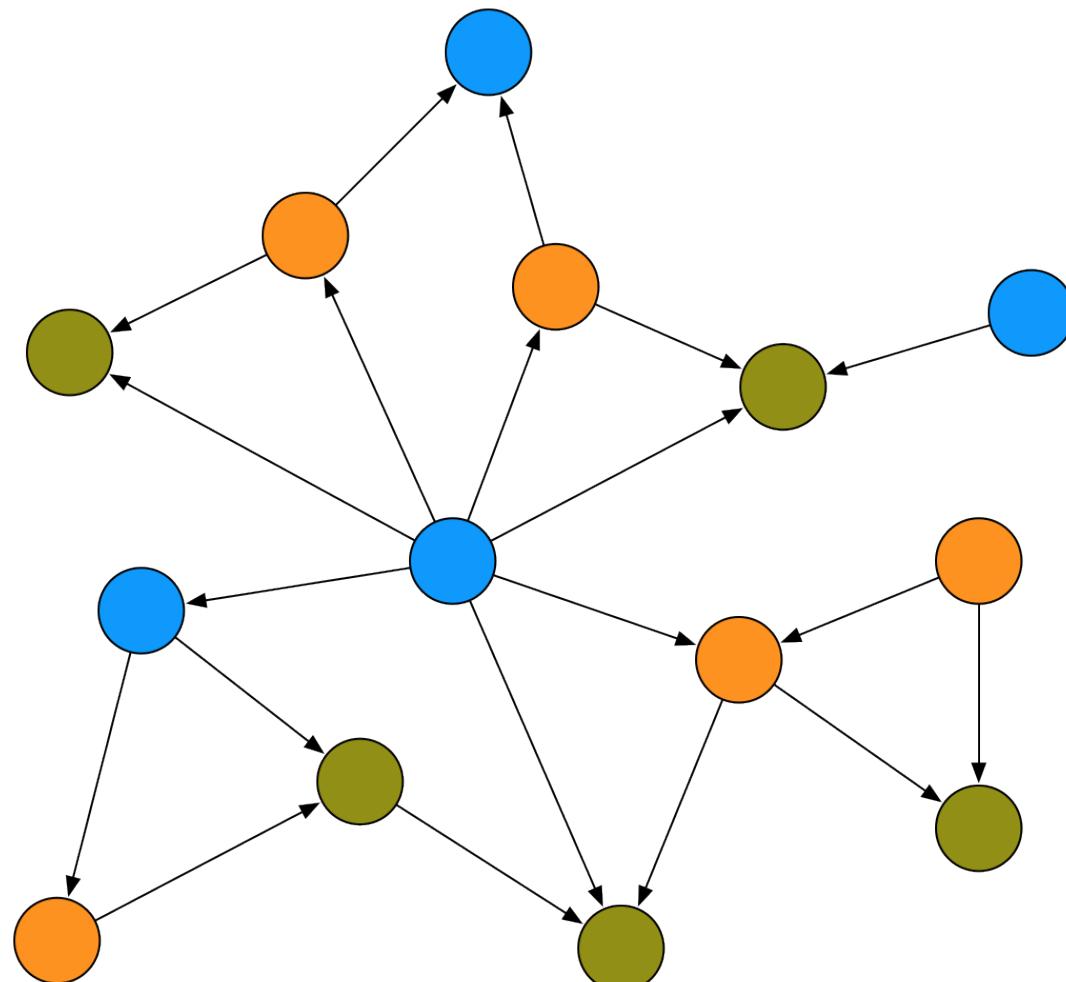


#neo4j

Pattern Matching

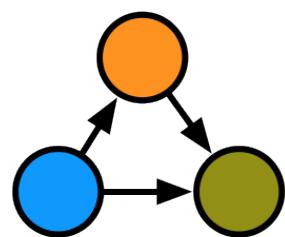


Pattern

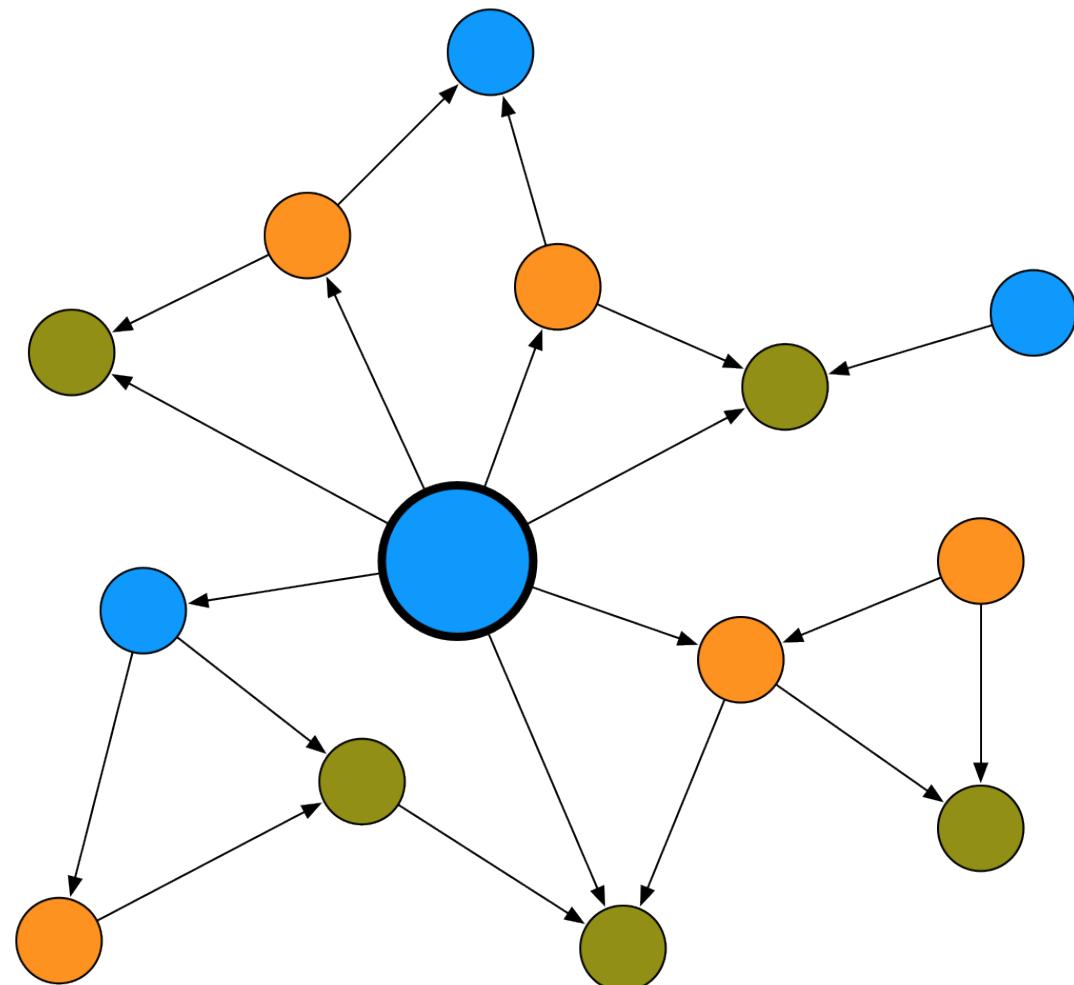


#neo4j

Start Node

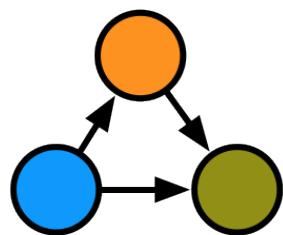


Pattern

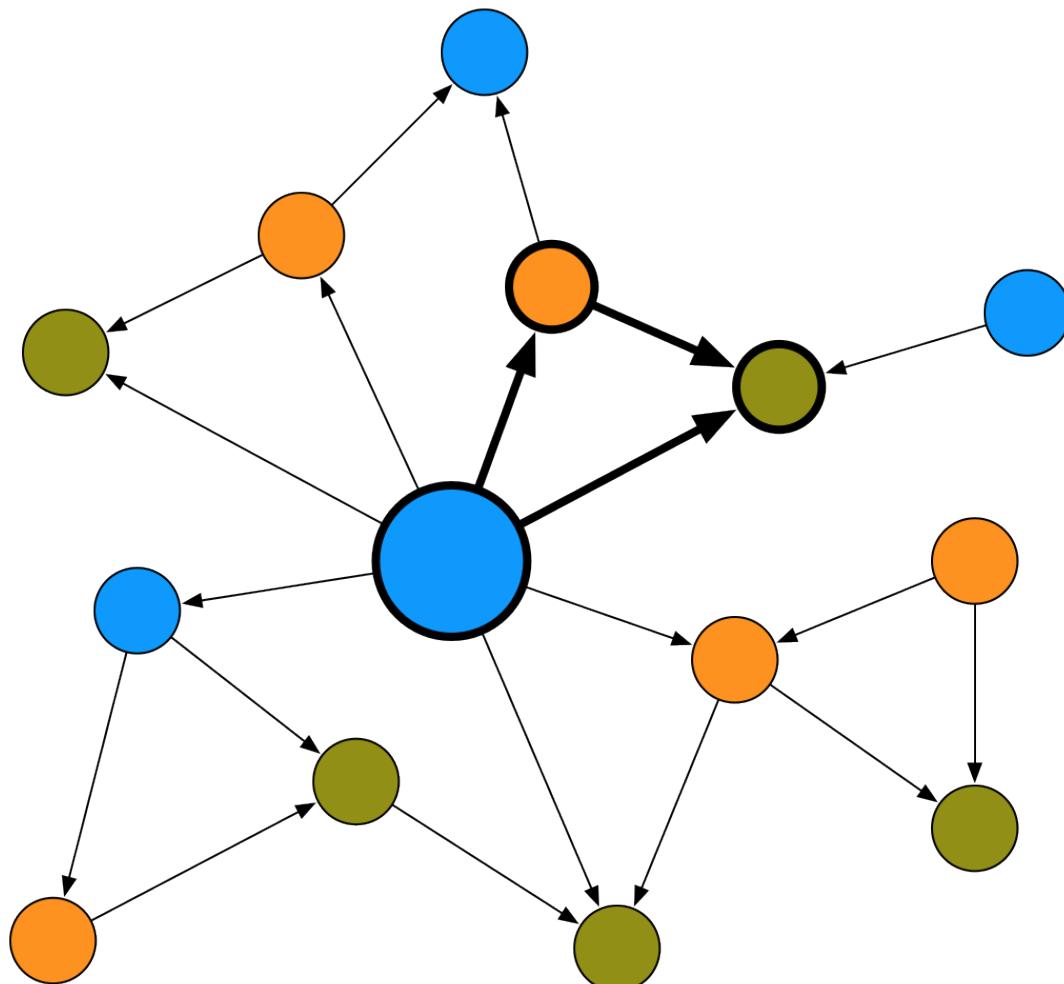


#neo4j

Match

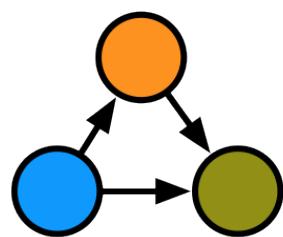


Pattern

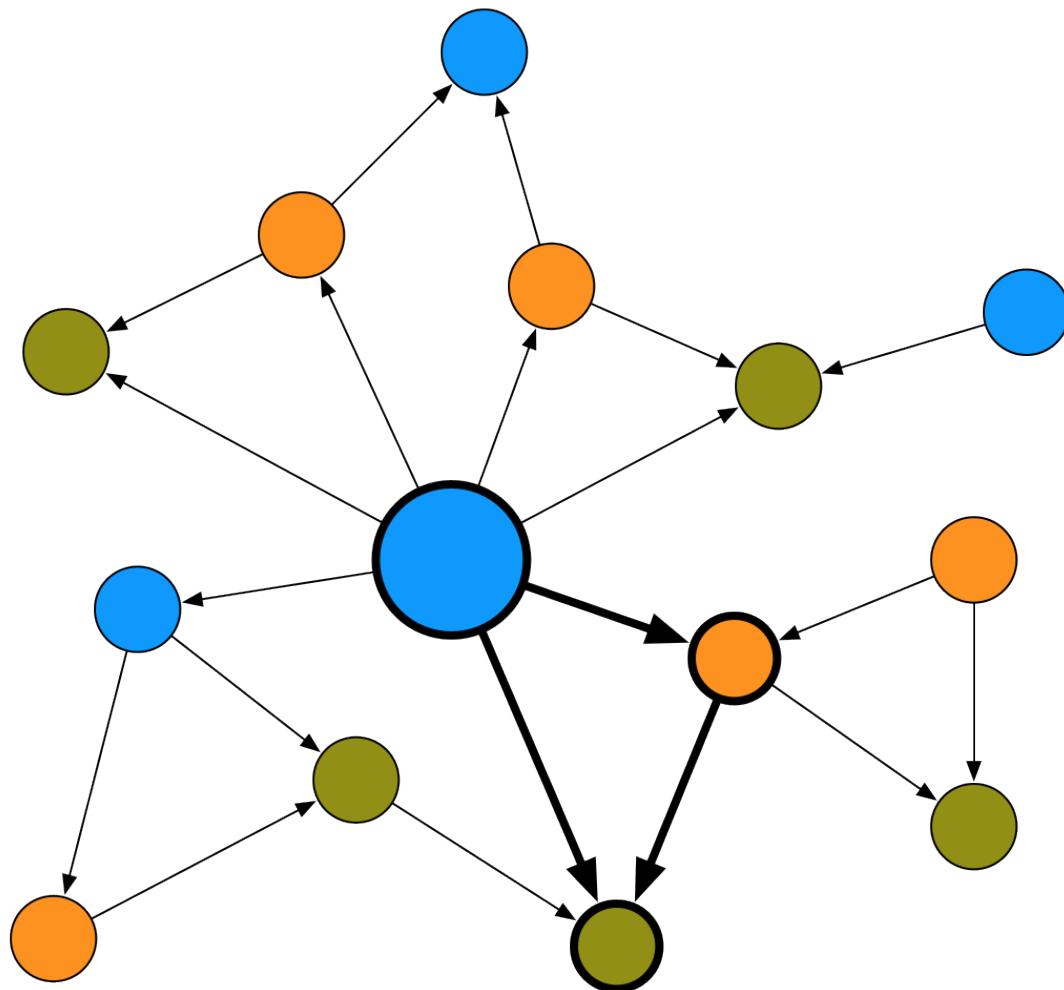


#neo4j

Match

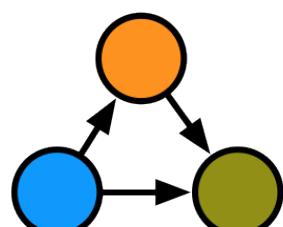


Pattern

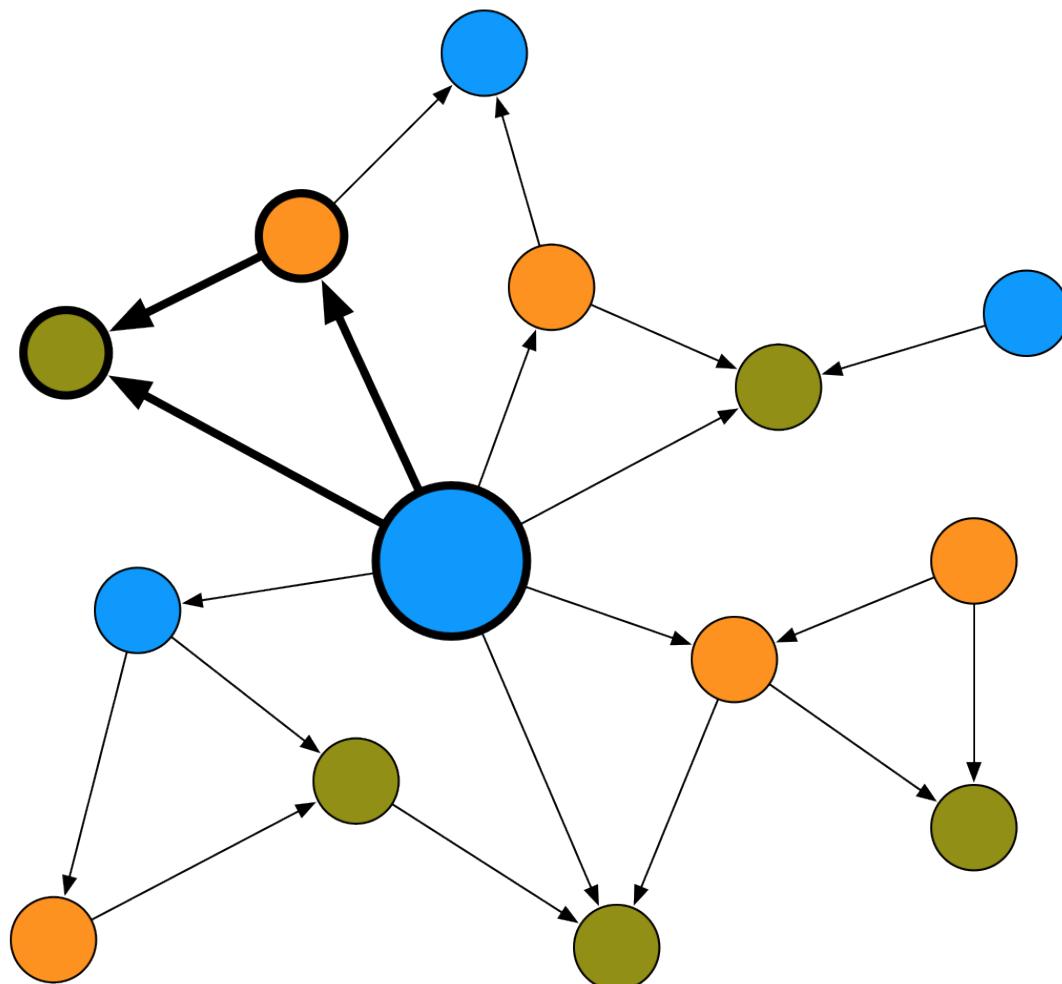


#neo4j

Match

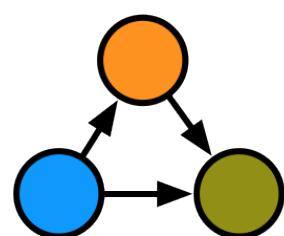


Pattern

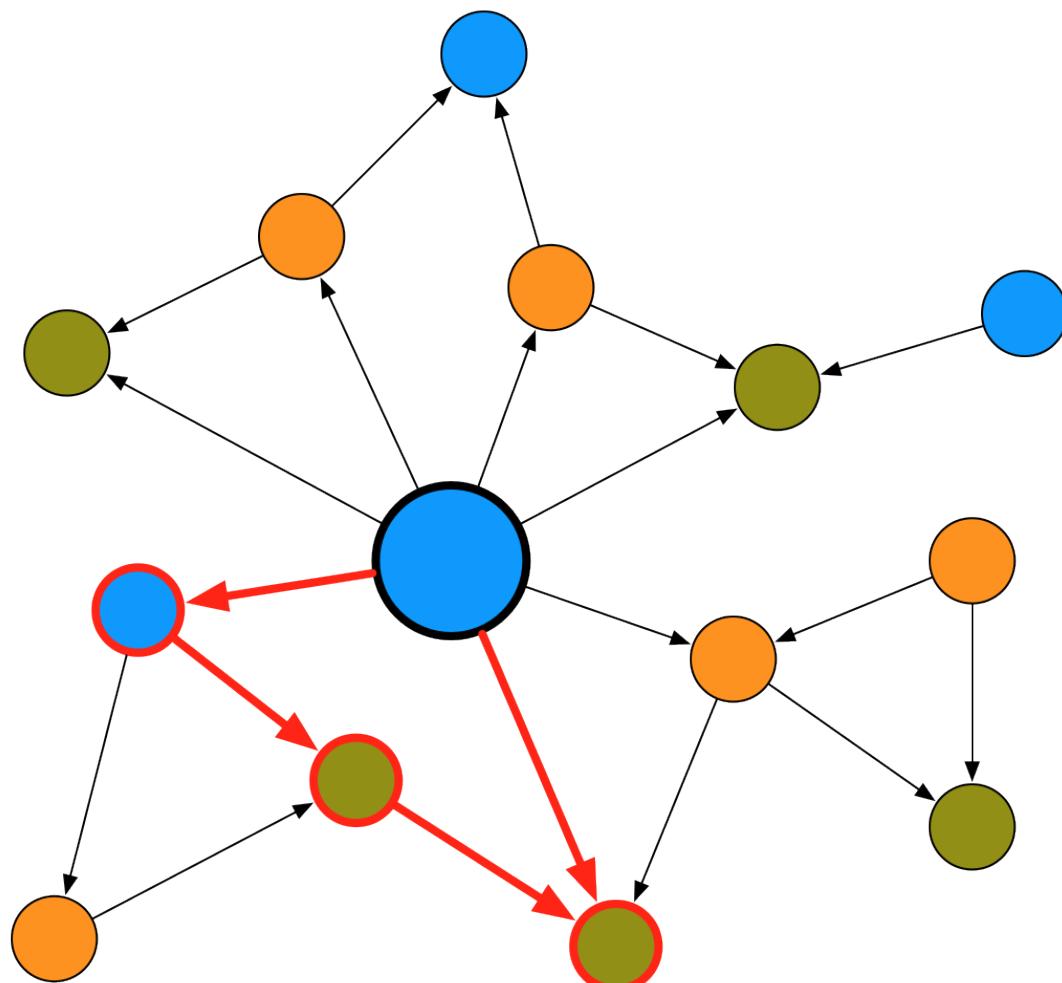


#neo4j

Non-Match

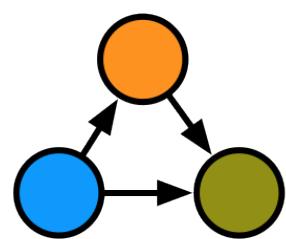


Pattern

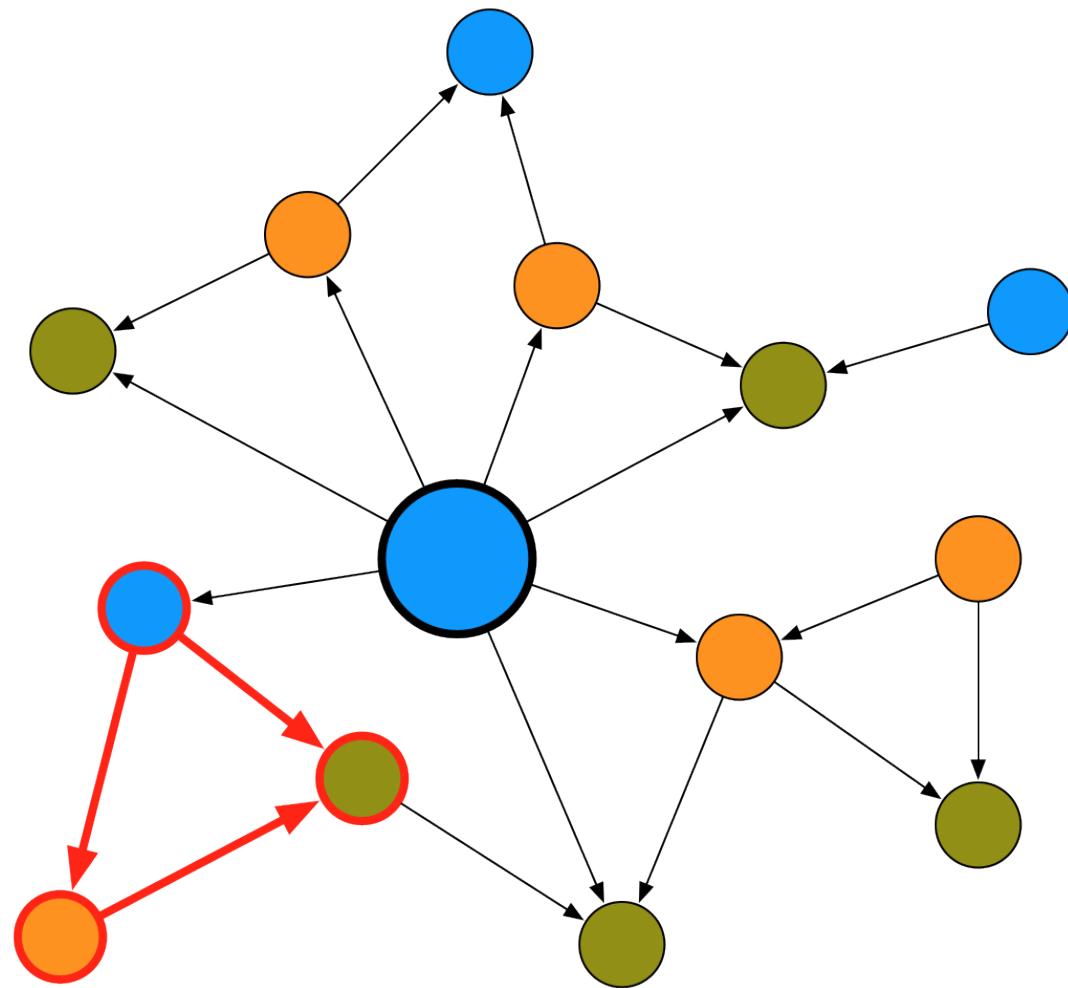


#neo4j

Non-Match

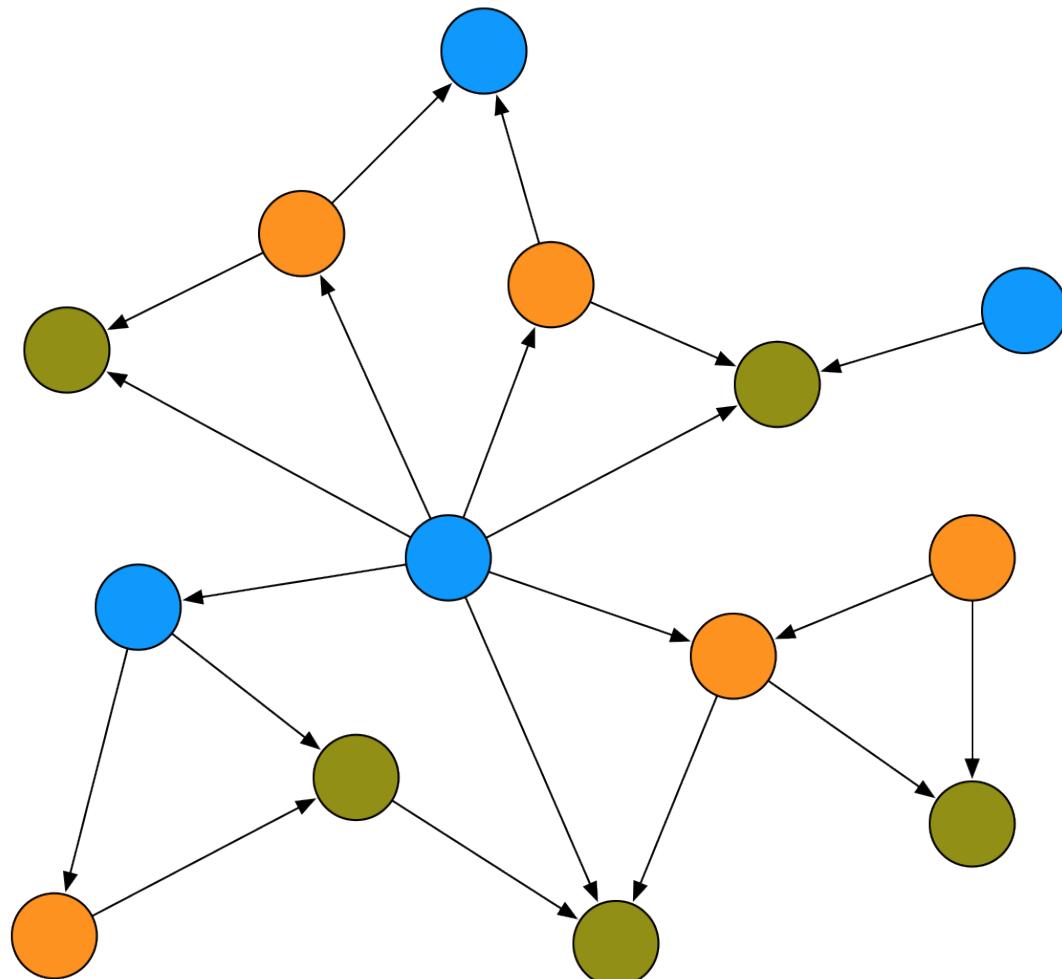


Pattern



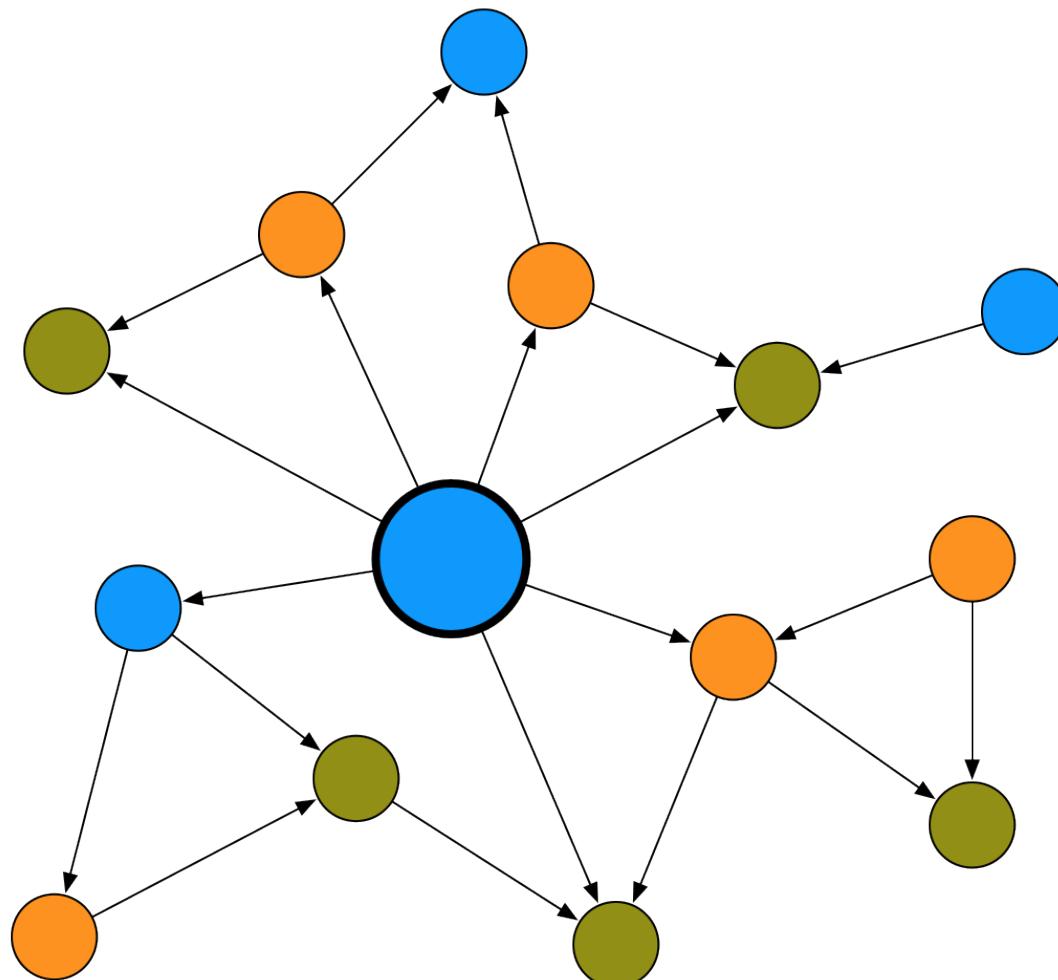
#neo4j

Traversal



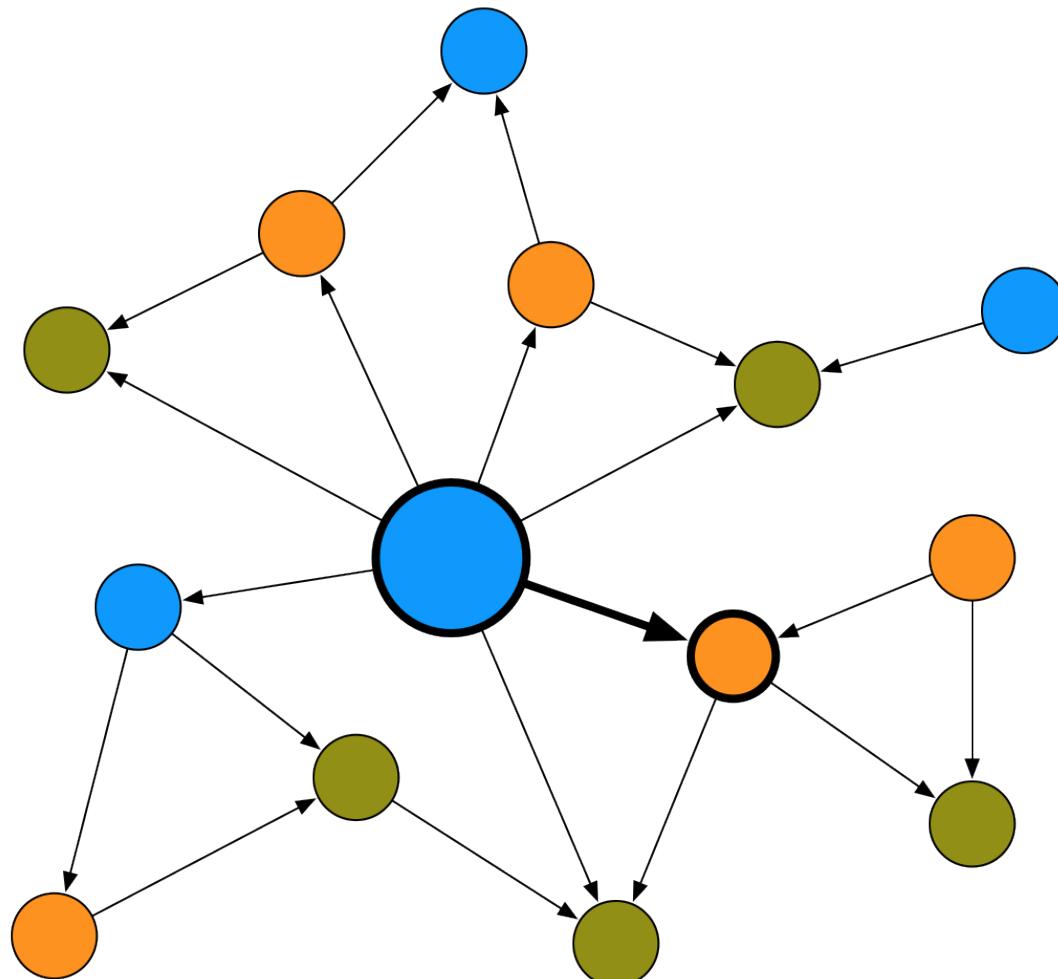
#neo4j

Start Node



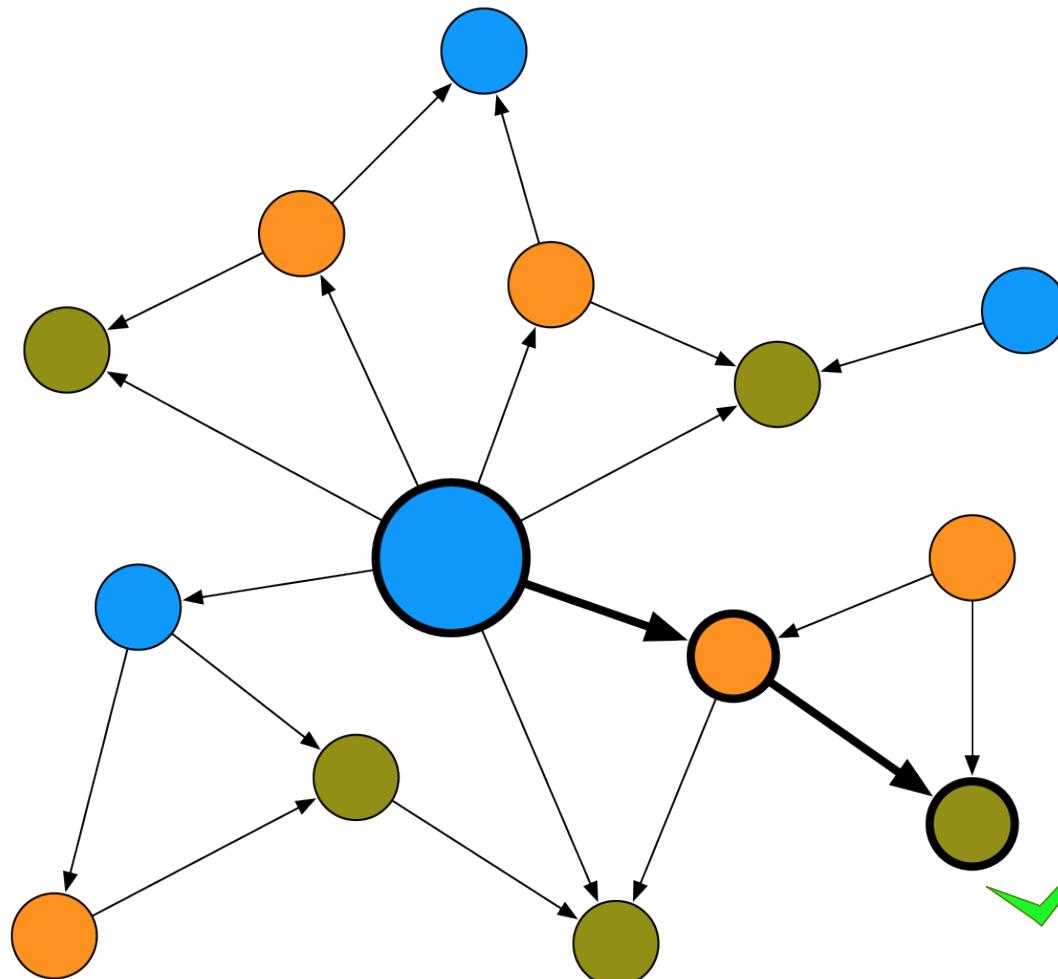
#neo4j

Follow Relationships



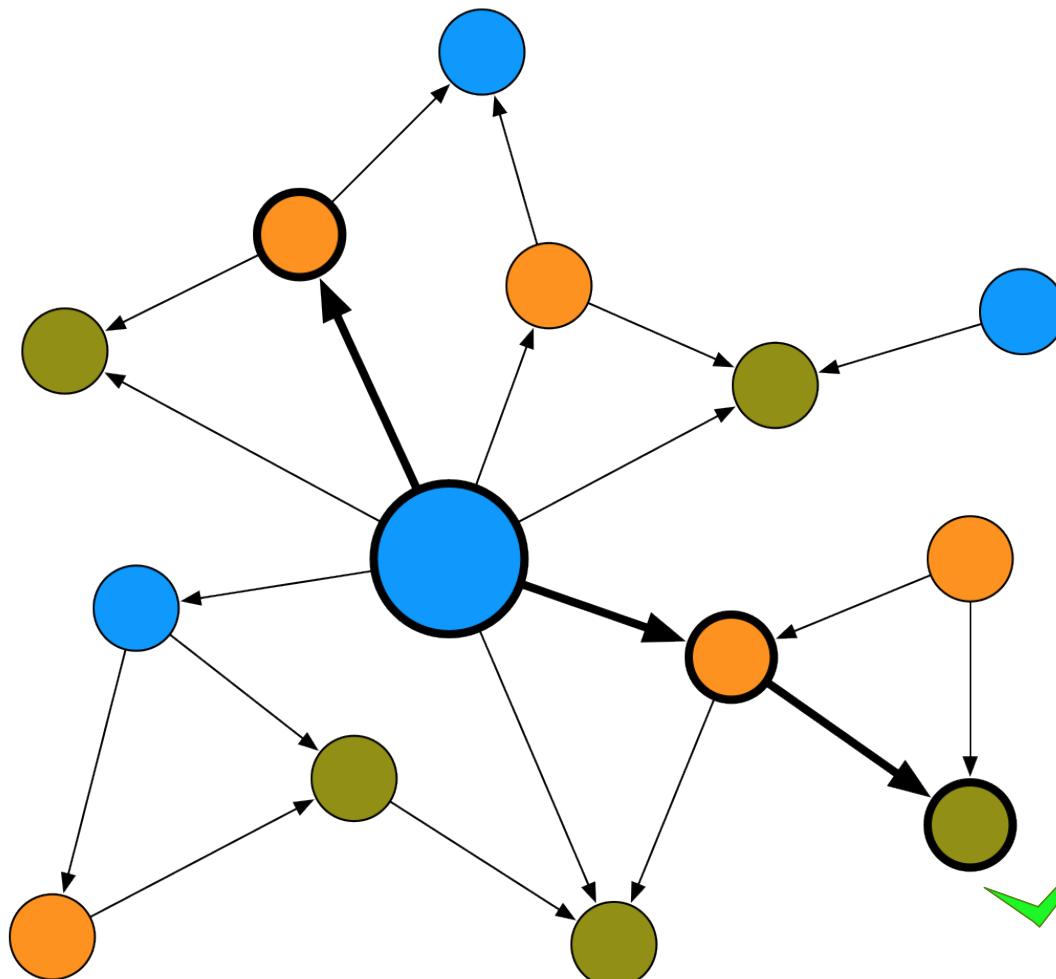
#neo4j

Evaluate Node



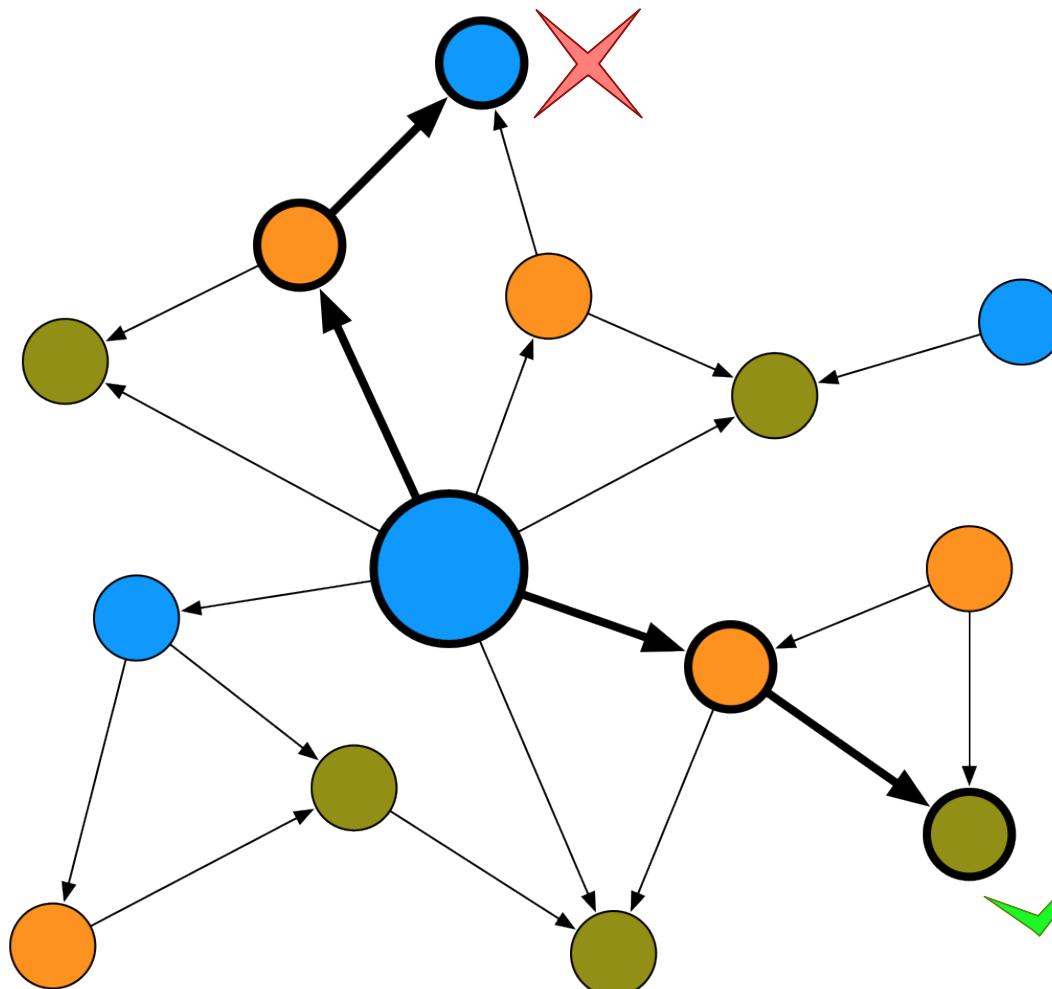
#neo4j

Continue Traversing



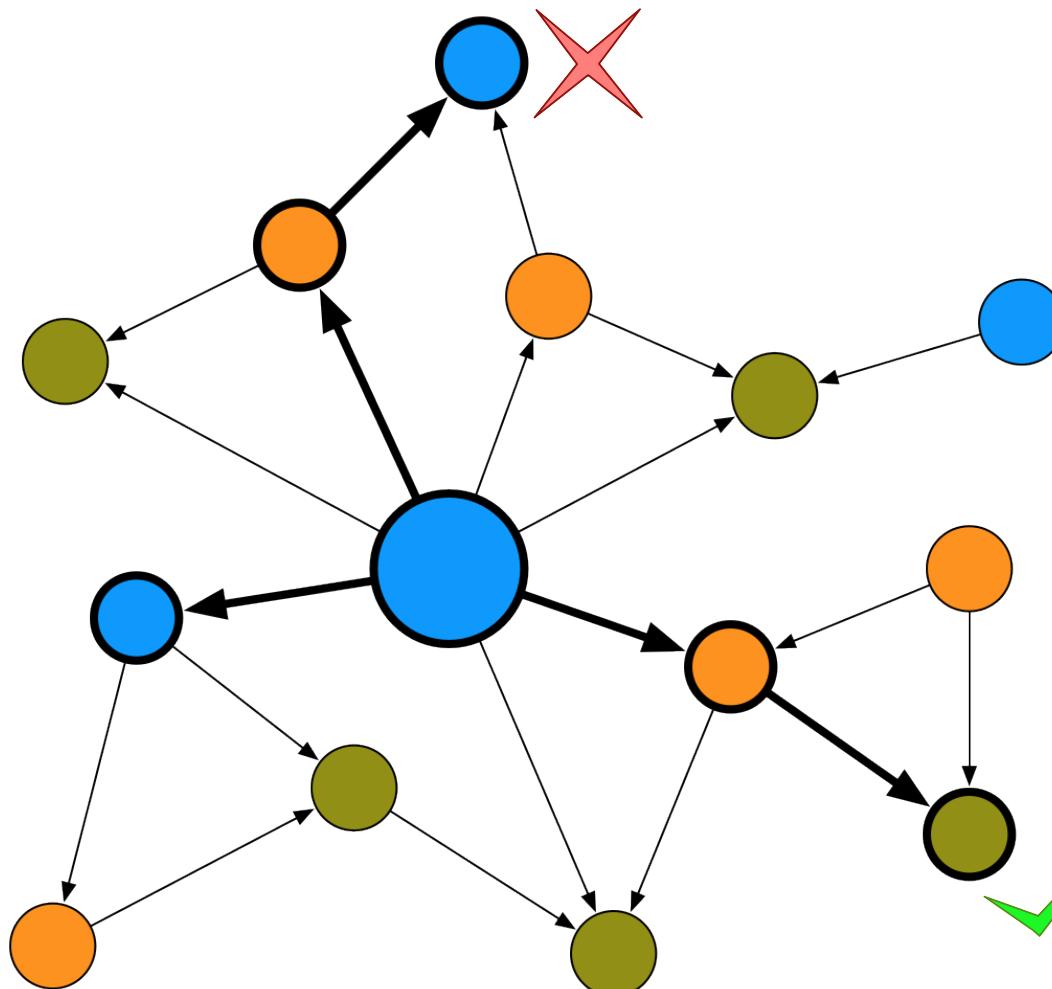
#neo4j

Continue Traversing



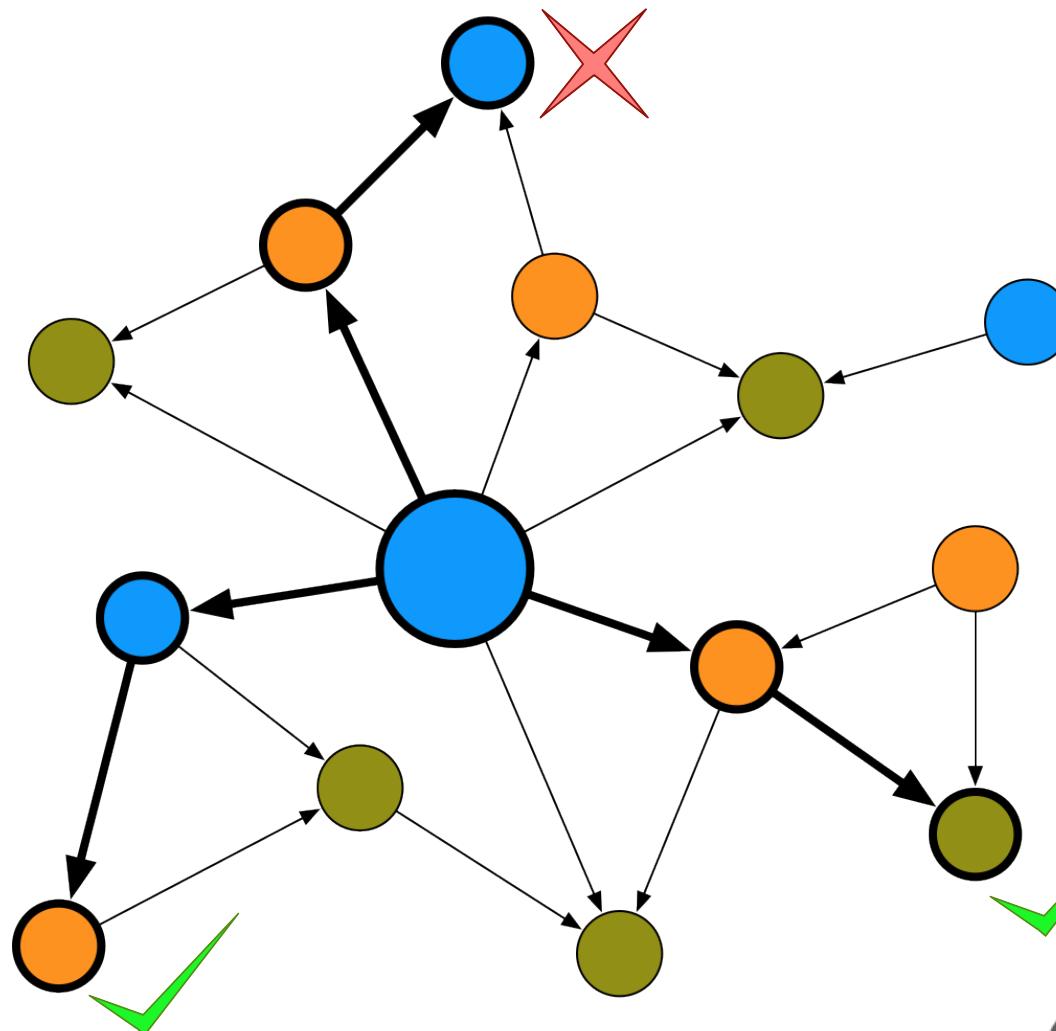
#neo4j

Continue Traversing



#neo4j

Continue Traversing



#neo4j

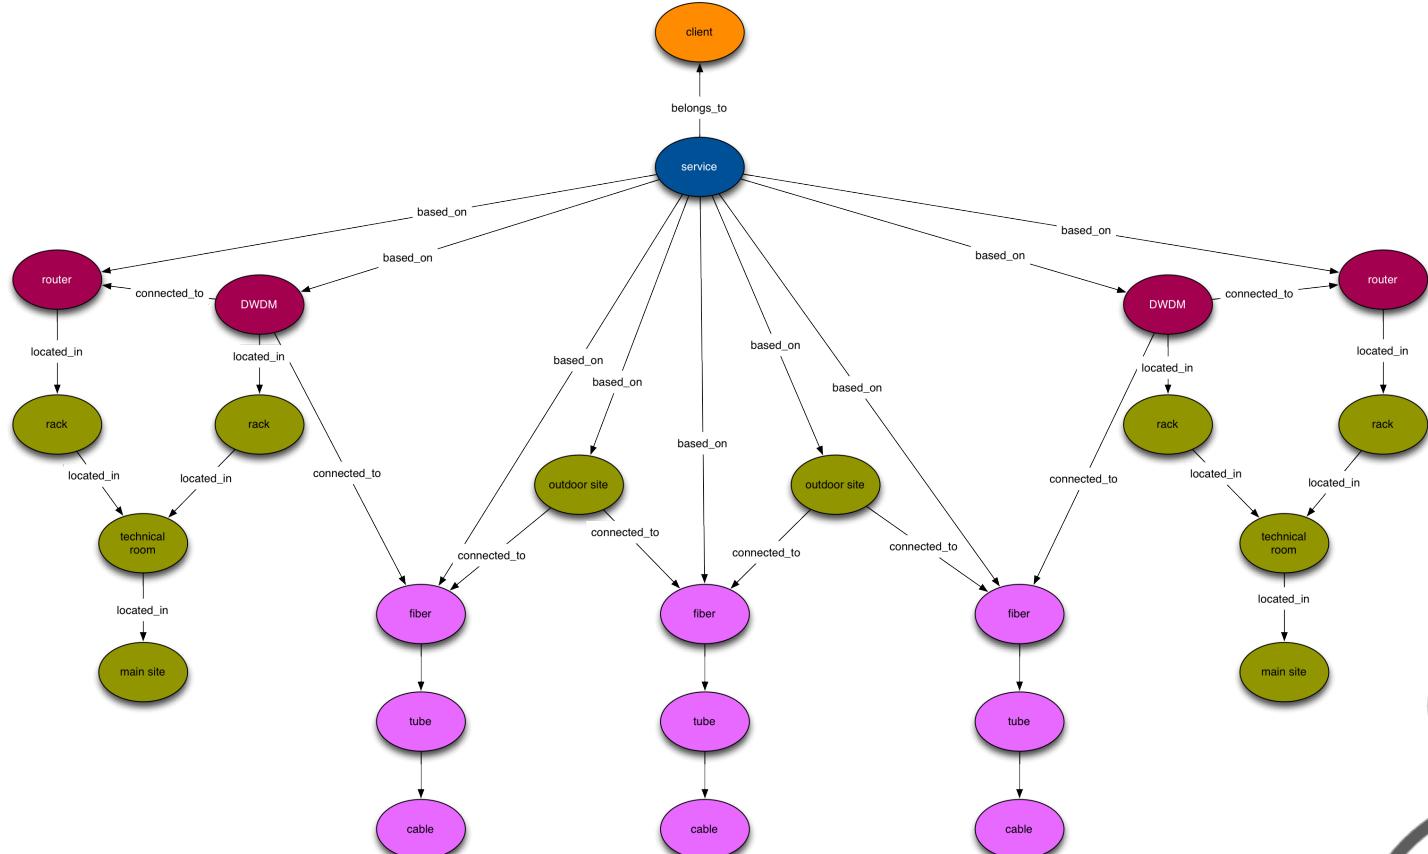
All Trades for Cost Centre

```
START cc=node:cost_centre(name='Cost Centre 1')  
MATCH cc-[*3]->trade  
RETURN trade;
```

Is Party Connected?

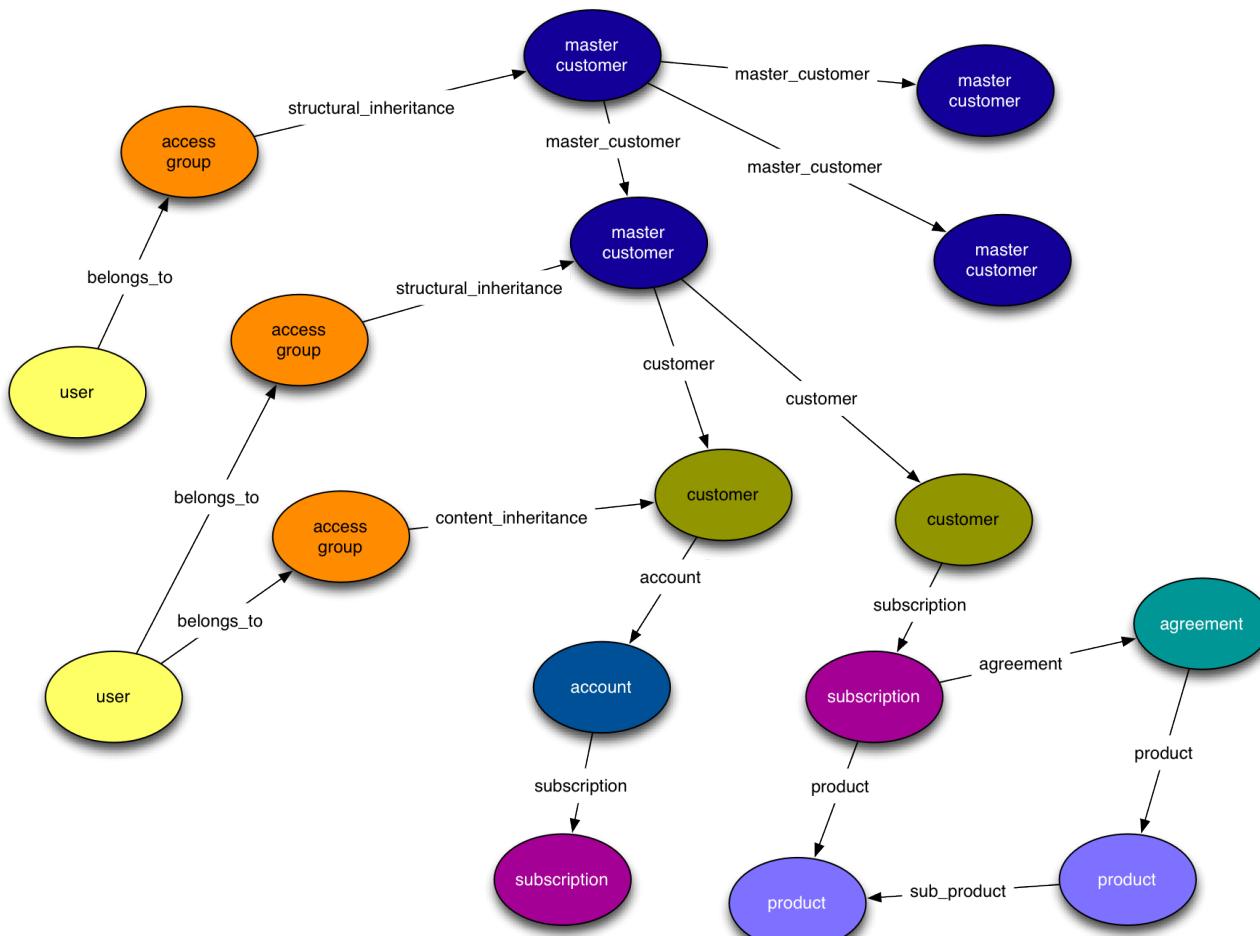
```
START cc=node:cost_centre(name='Cost Centre 1'),  
     party=node:party(name='Party 1')  
MATCH p=shortestPath(cc-[ *..5]-party)  
RETURN length(p);
```

Network Impact Analysis



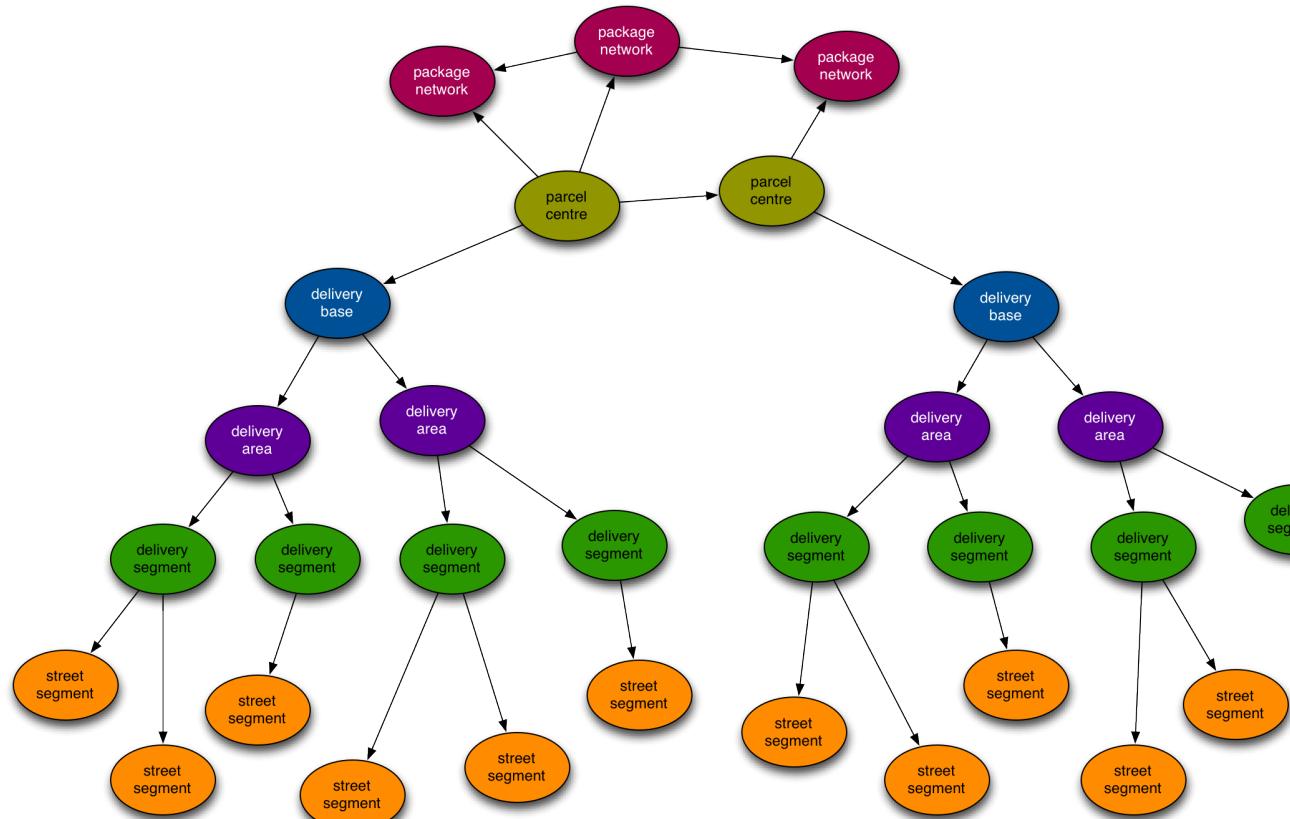
#neo4j

Asset Management & Access Control



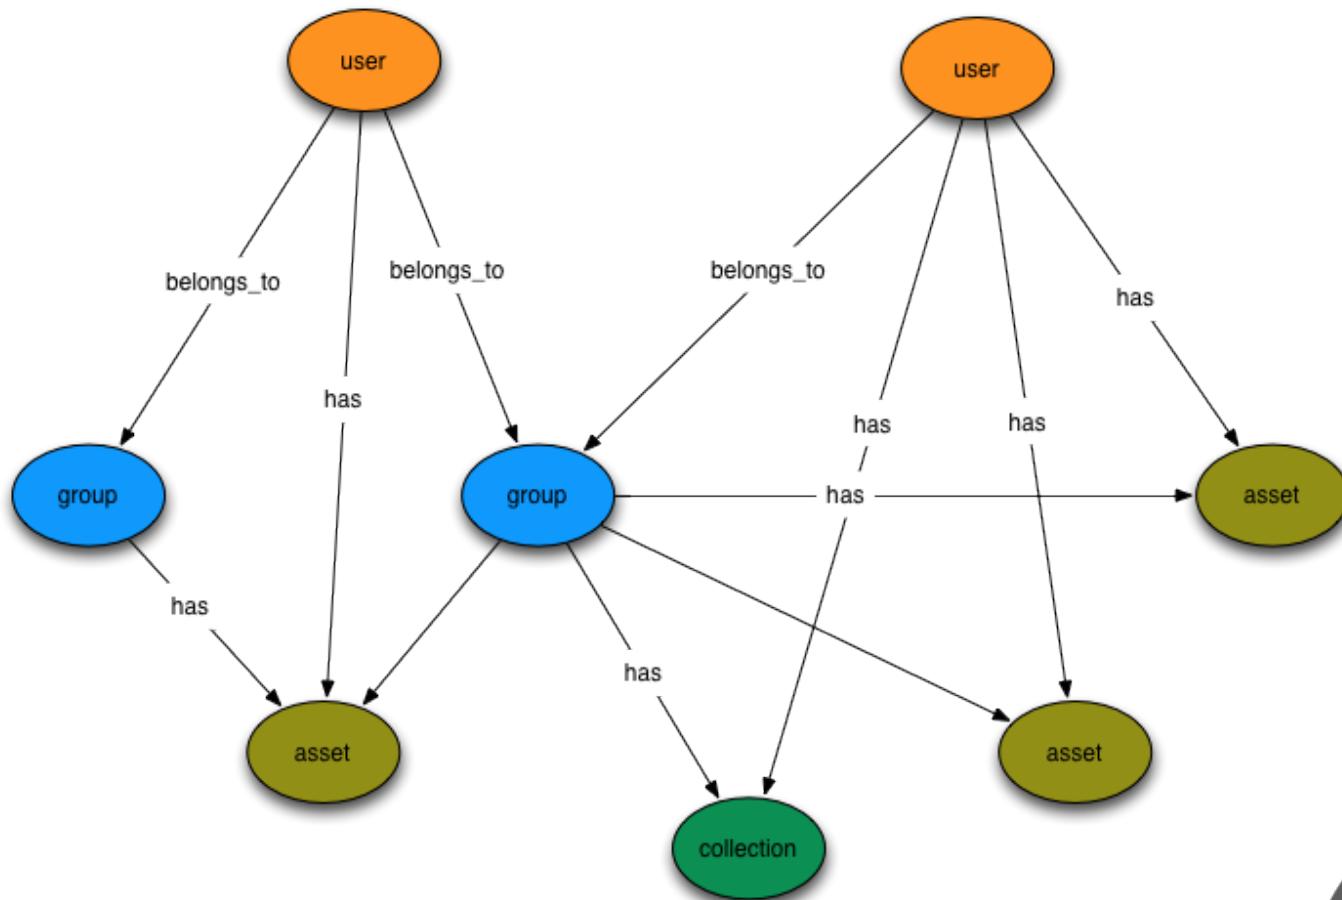
#neo4j

Logistics



#neo4j

Social Network & Recommendations

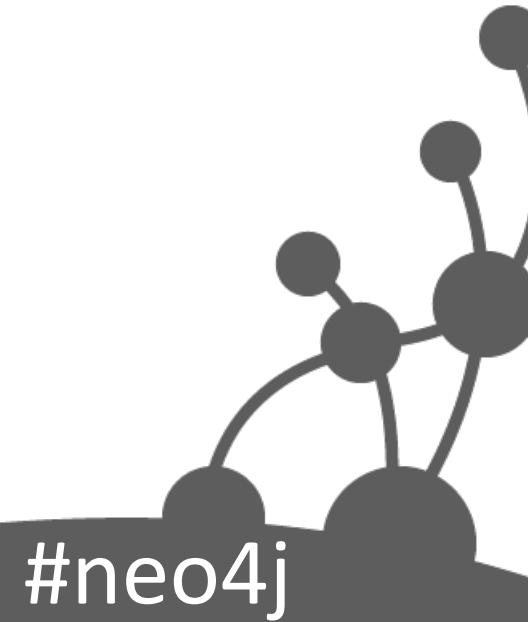


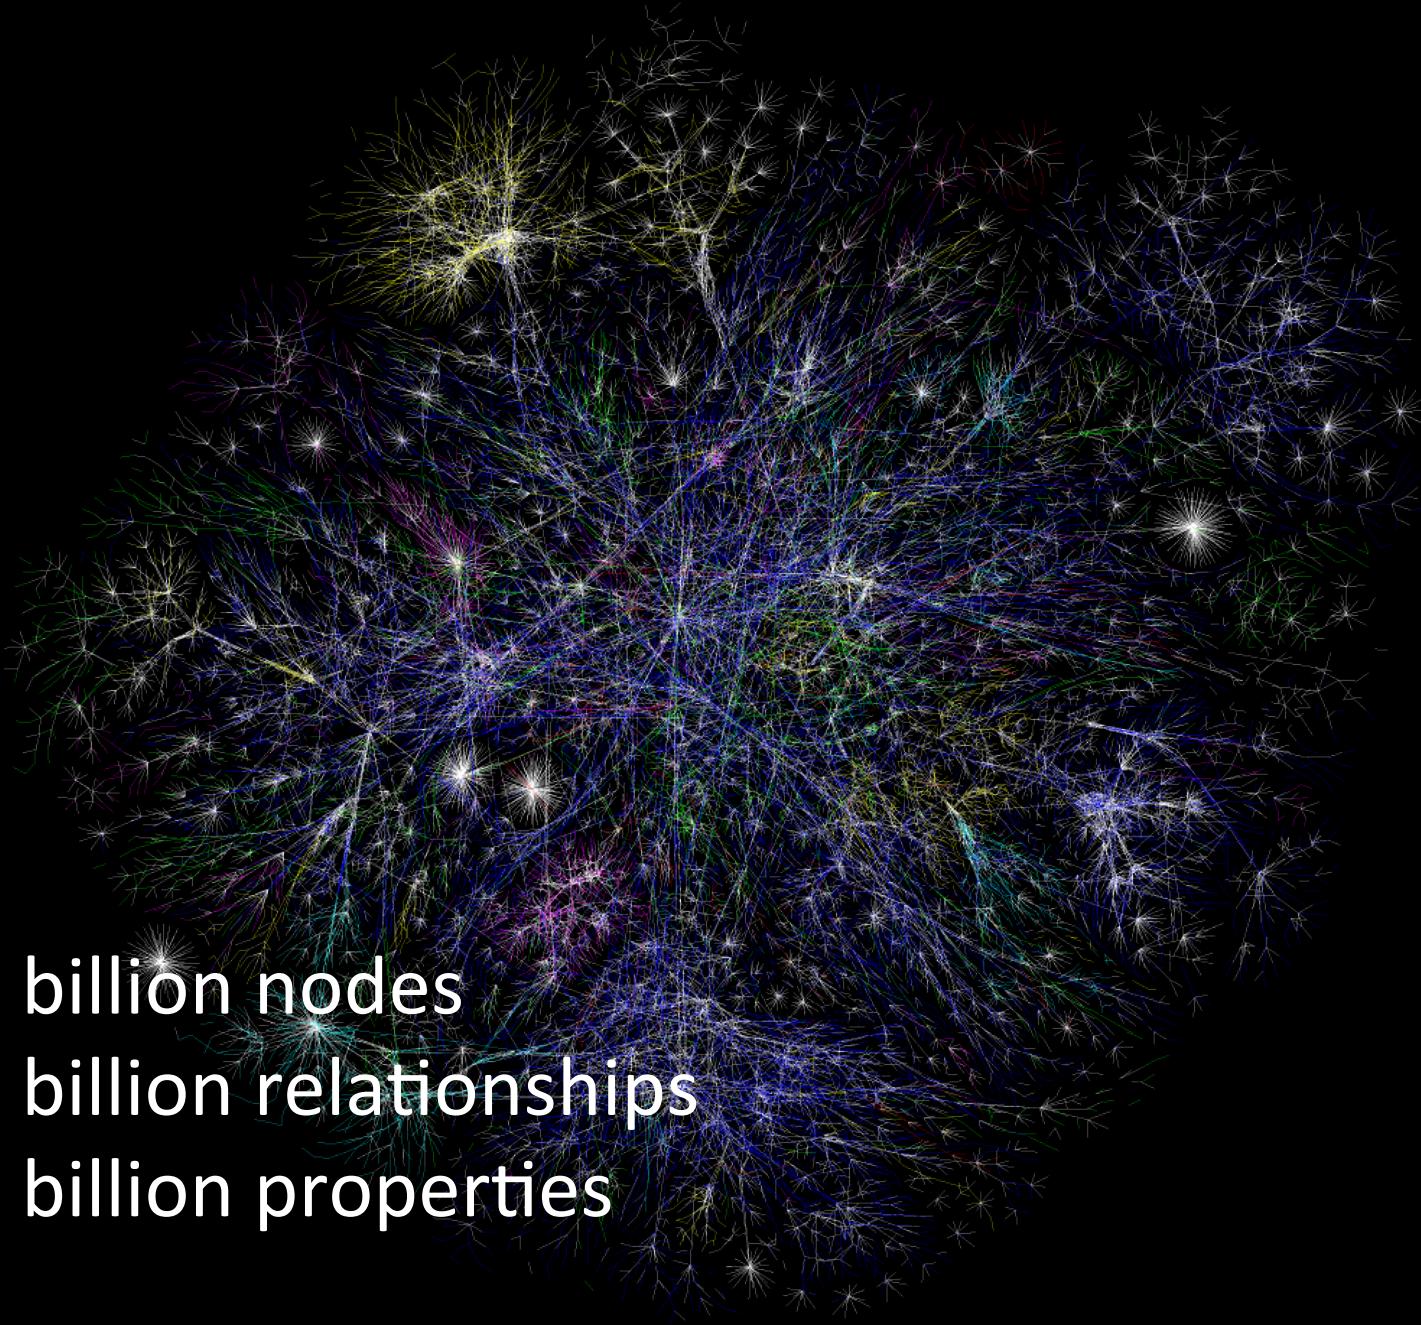
#neo4j



Questions

@iansrobinson
ian.robinson@neotechnology.com





A complex network graph visualization featuring a dense cluster of nodes and a multitude of connecting edges. The nodes are represented by small dots, and the edges are thin lines of various colors, including yellow, green, blue, and red. The graph is set against a black background and has a radial, star-like structure with many branches extending from a central point.

32 billion nodes
32 billion relationships
64 billion properties